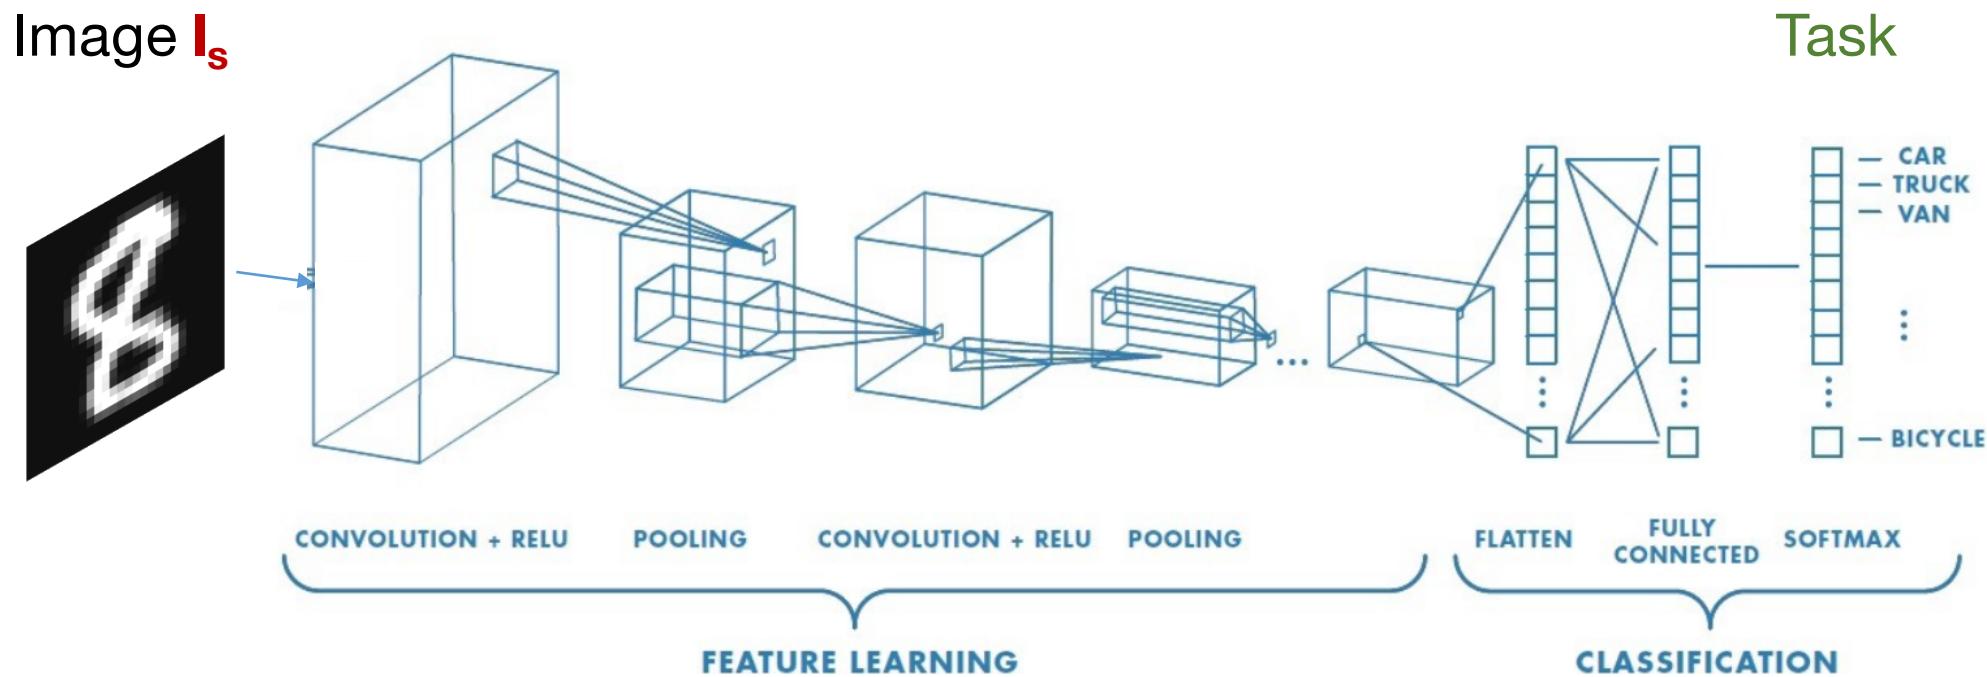


Machine Learning in Imaging

BME 590L
Roarke Horstmeyer

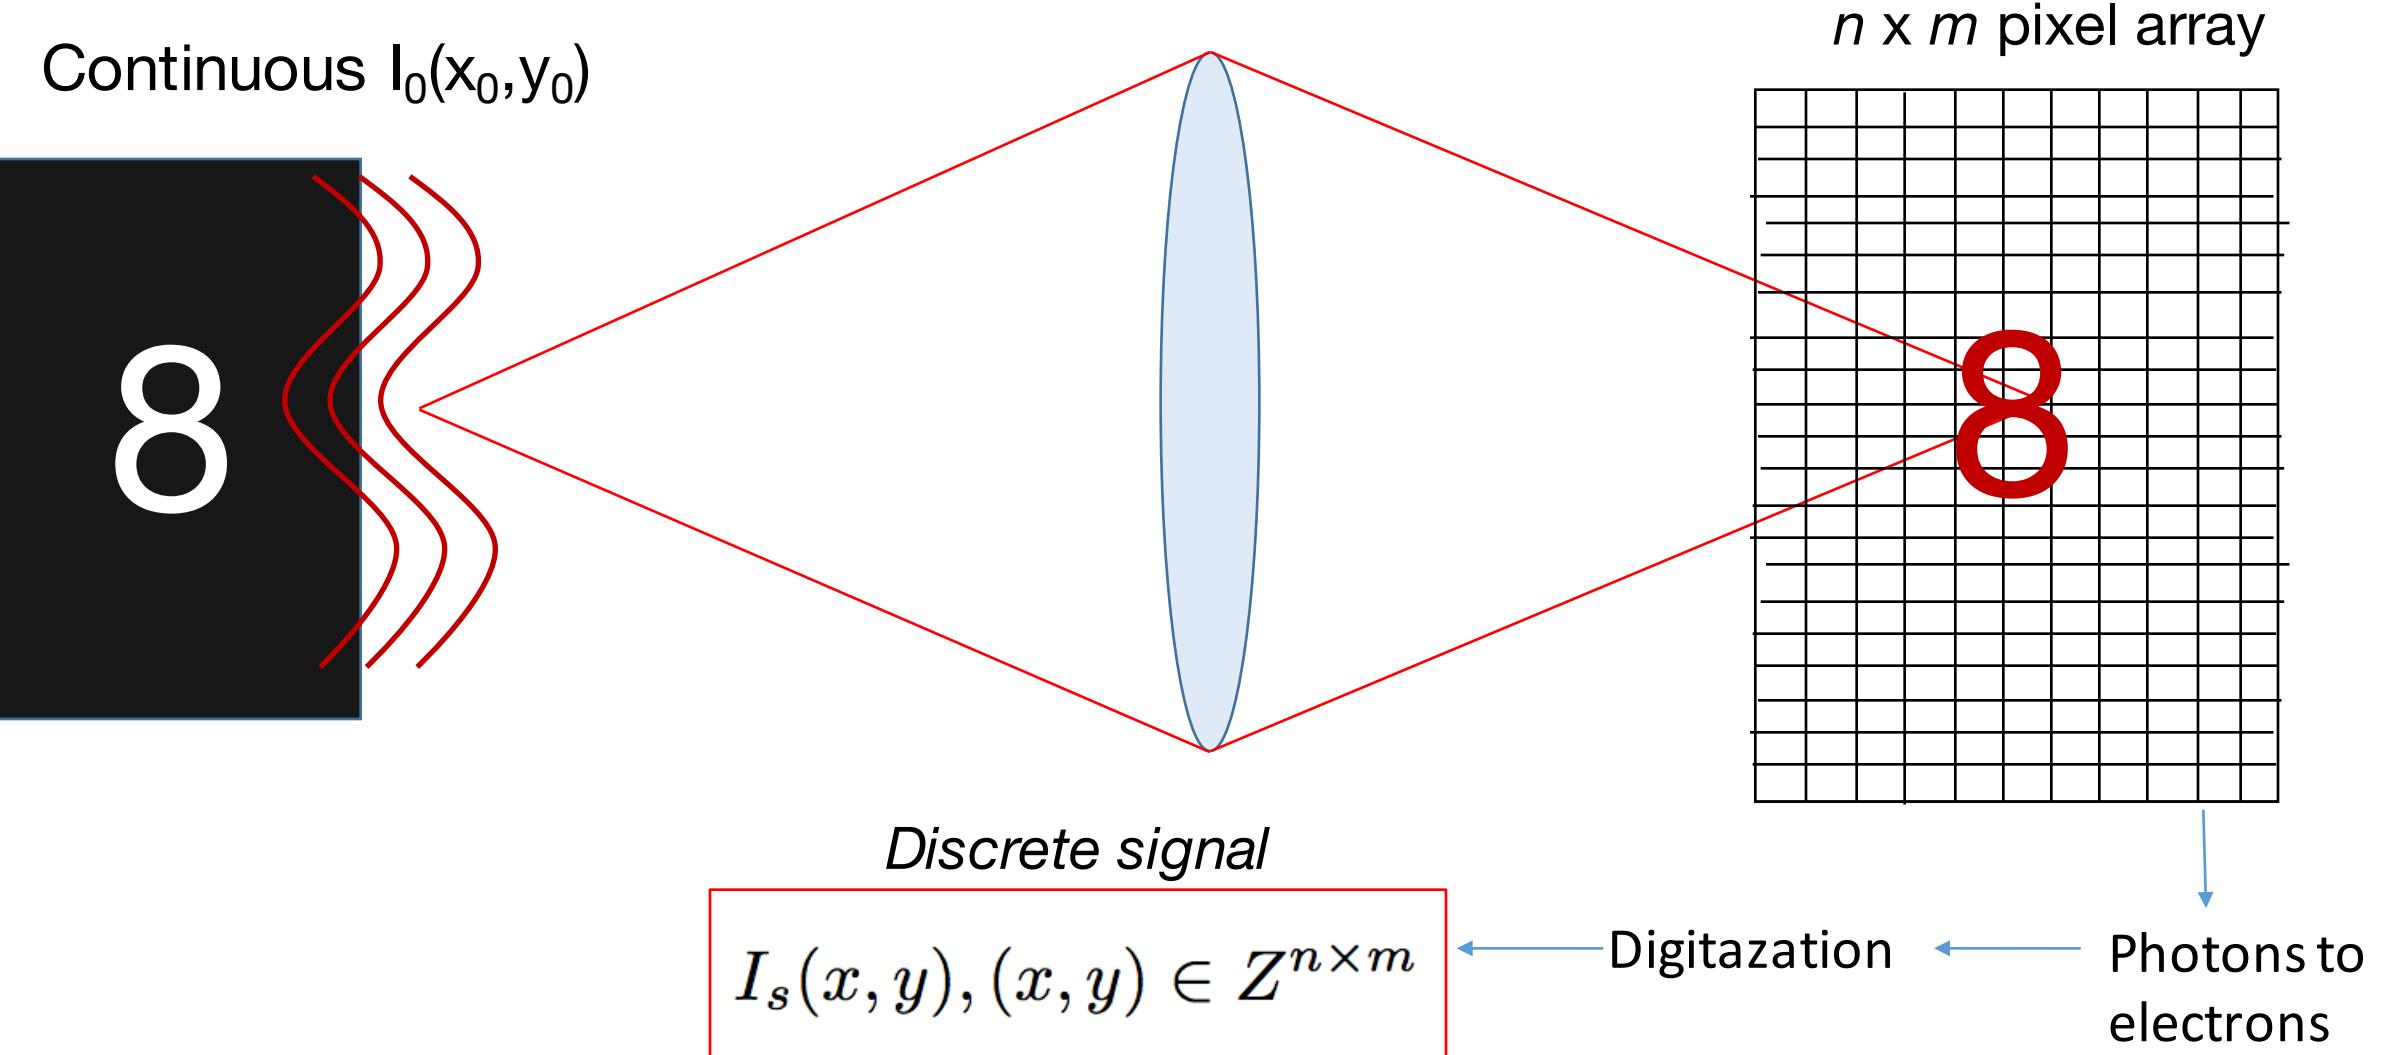
Lecture 16: Physical CNN examples & introduction to optics

Bringing together physical and digital image representations



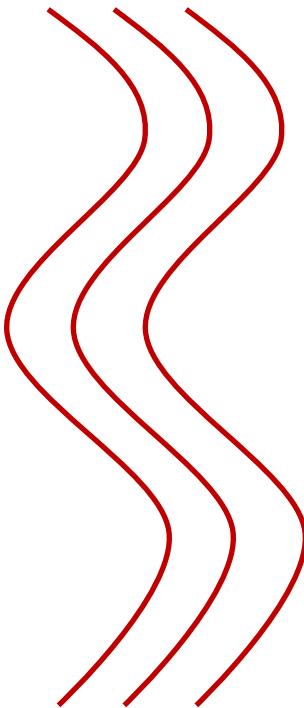
$$\text{Task} = W_n \dots \text{ReLU}[W_1 \text{ ReLU}[W_0 I_s] \dots]$$

Simple model of image formation



What does the Sampling Theorem mean for us?

Continuous functions



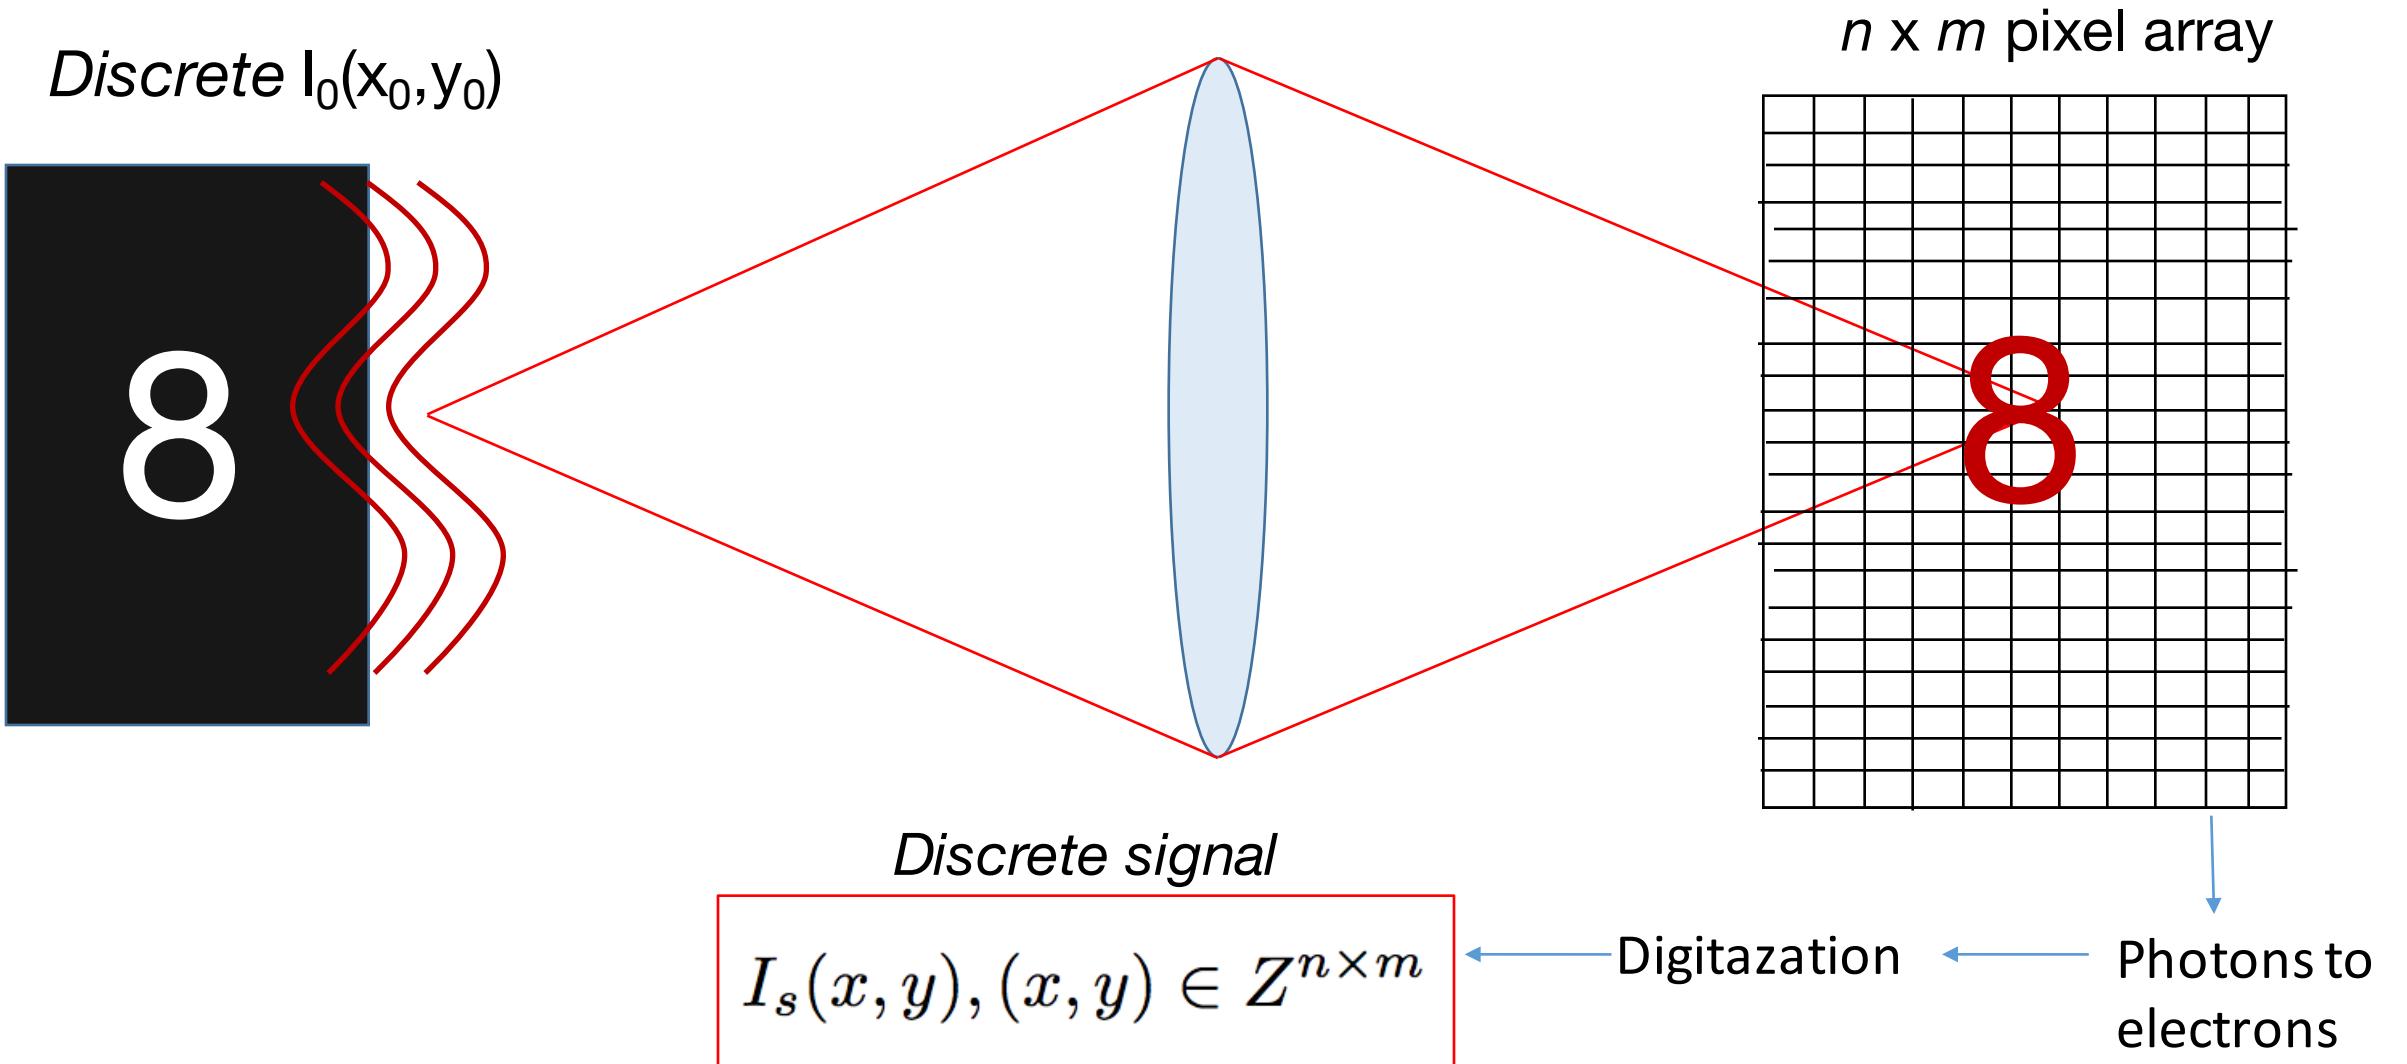
*conditions

Discretize vectors
(and matrices)

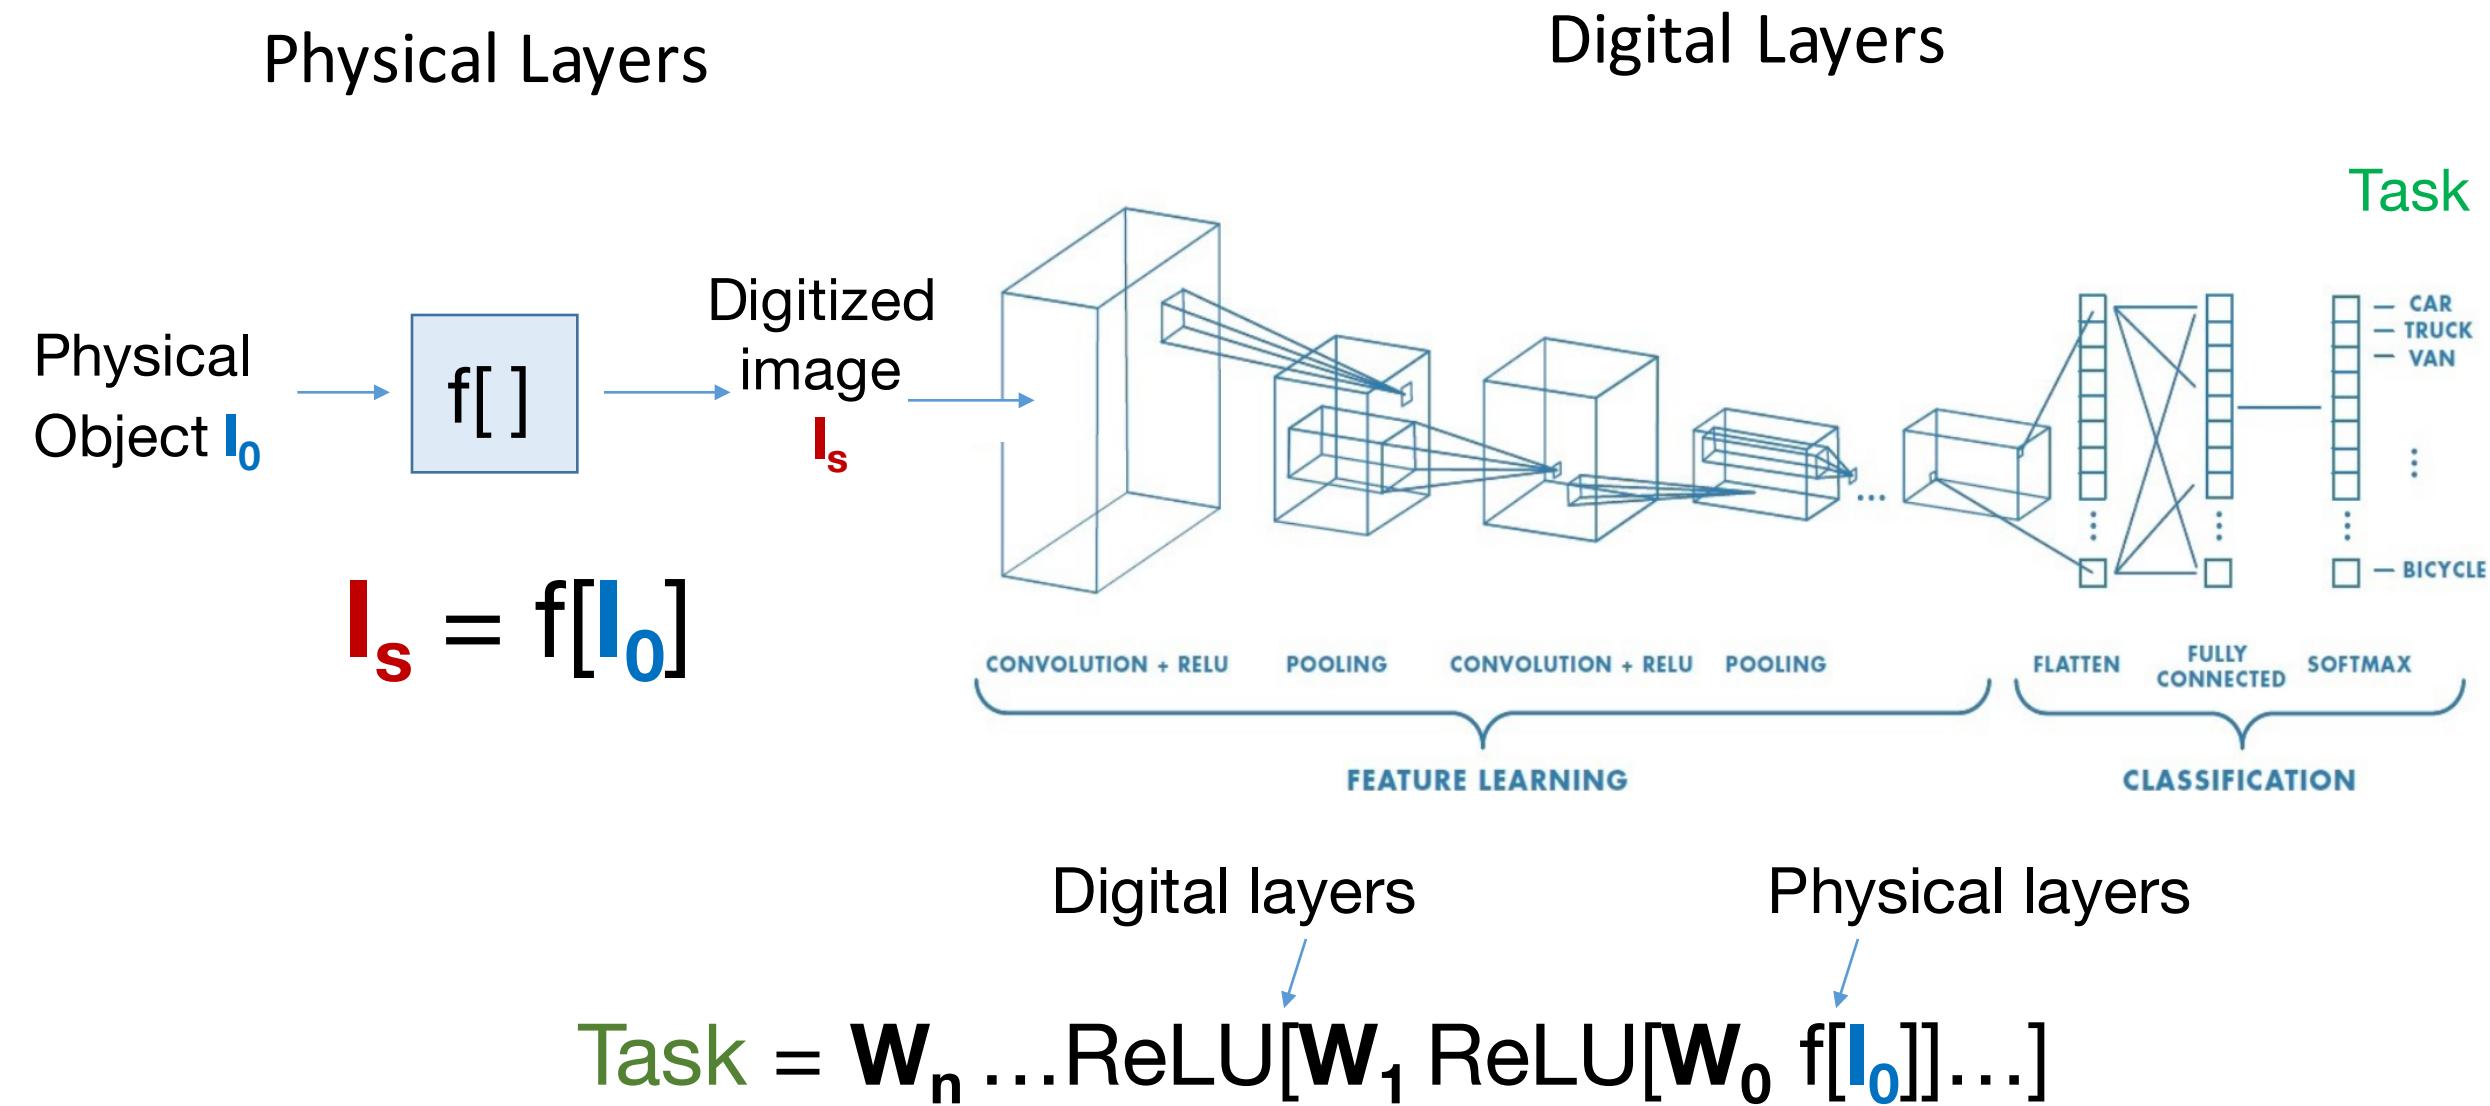


17
20
22
21
23
25
24
26
29

Simple model of image formation

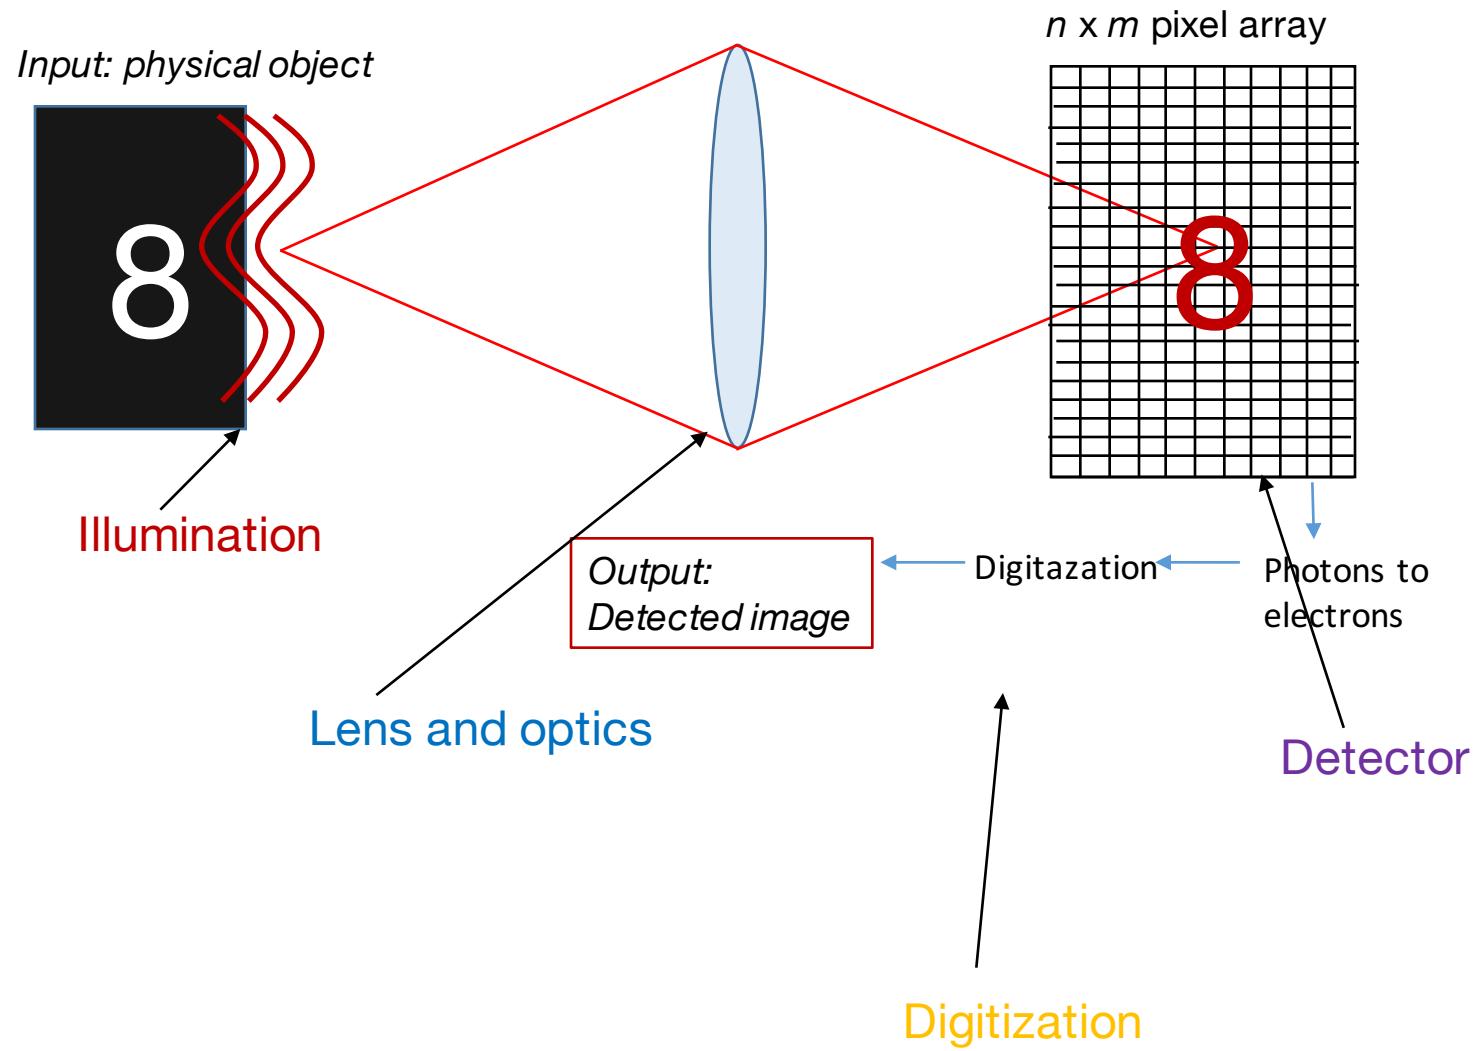


Bringing together physical and digital image representations

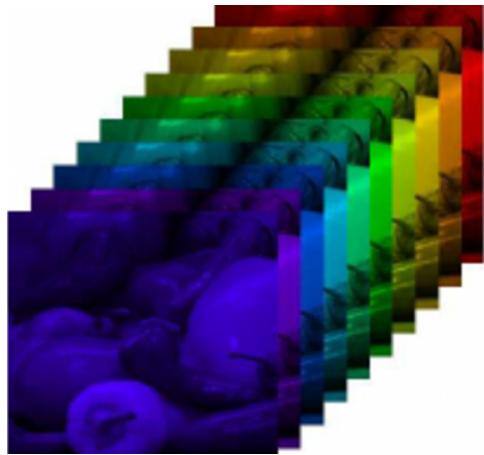


What physical parameters effect image formation $f[]$?

- **Illumination**
 - Spatial pattern
 - Angle of incidence
 - Color, polarization
- **Lens and optics**
 - Position/orientation
 - Shape
 - Focus
 - Transparency
- **Detector**
 - Pixel size
 - Pixel shape & fill factor
 - Color filters
 - Other filters
- **Digitization**
 - E to P curves
 - Digitization schemes/thresholds
 - Data transmission, multiplexing
- Physical object



Example #1: Optimized color filter for a grayscale camera

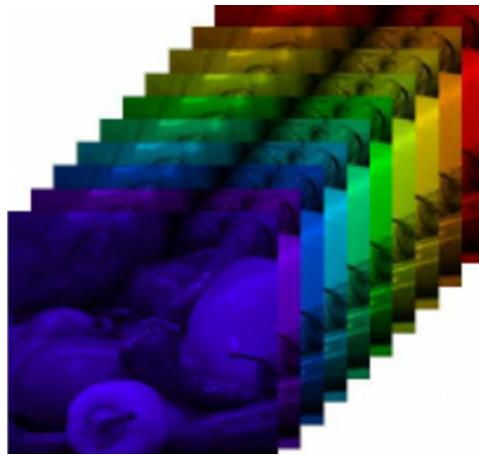


“Ground truth” object:

$$I_0(x, y, \lambda)$$

100 x 100 pix. x 30 spectral channels

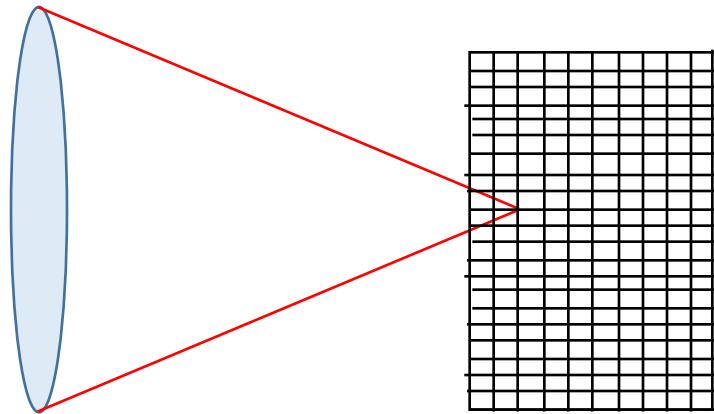
Example #1: Optimized color filter for a grayscale camera



“Ground truth” object:

$$I_0(x, y, \lambda)$$

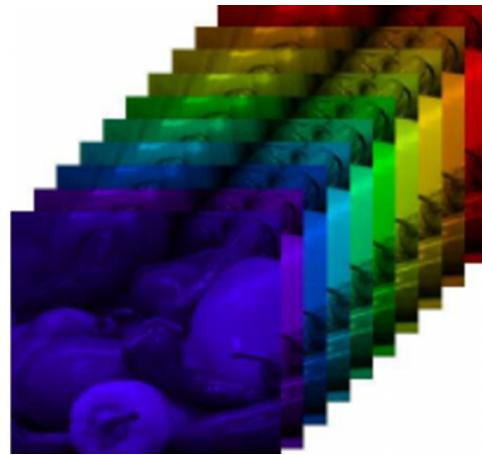
100 x 100 pix. x 30 spectral channels



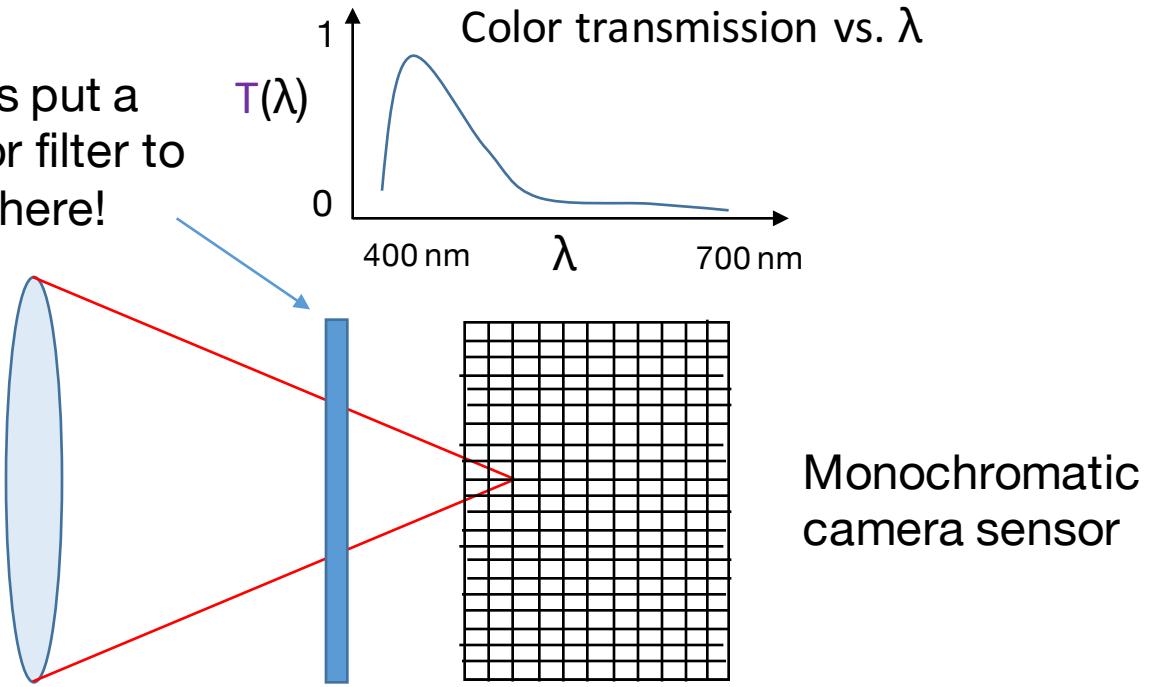
Monochromatic
camera sensor

$$I_s(x, y) = \sum_{\lambda} I_0(x, y, \lambda)$$

Example #1: Optimized color filter for a grayscale camera



Let's put a
color filter to
put here!



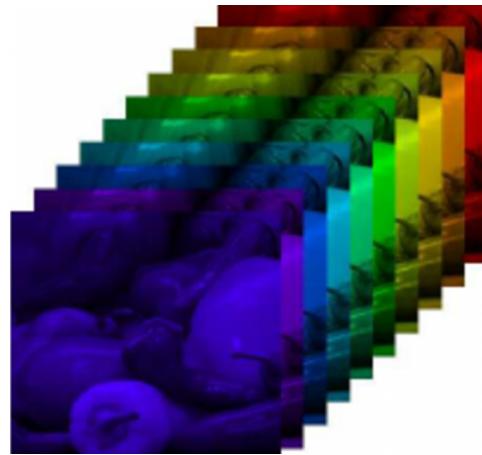
“Ground truth” object:

$$I_0(x, y, \lambda)$$

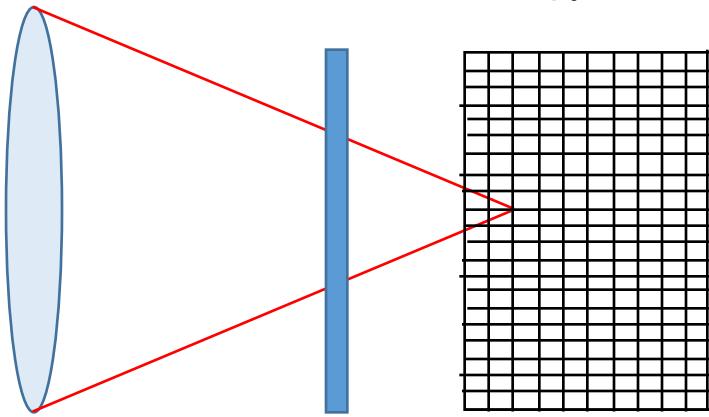
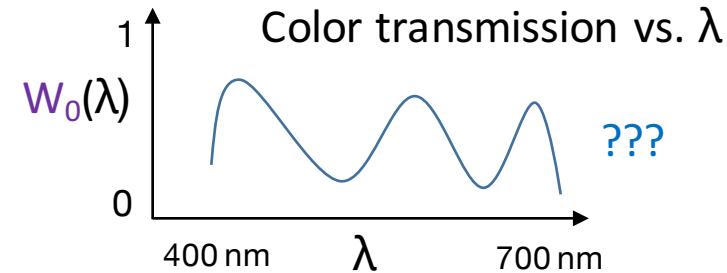
100 x 100 pix. x 30 spectral channels

$$I_s(x, y) = \sum_{\lambda} T(\lambda) I_0(x, y, \lambda)$$

Example #1: Optimized color filter for a grayscale camera



Design optimal
color filter for
classification:



Monochromatic
camera sensor

Training data:

$$[I_0(x, y, \lambda), y]$$

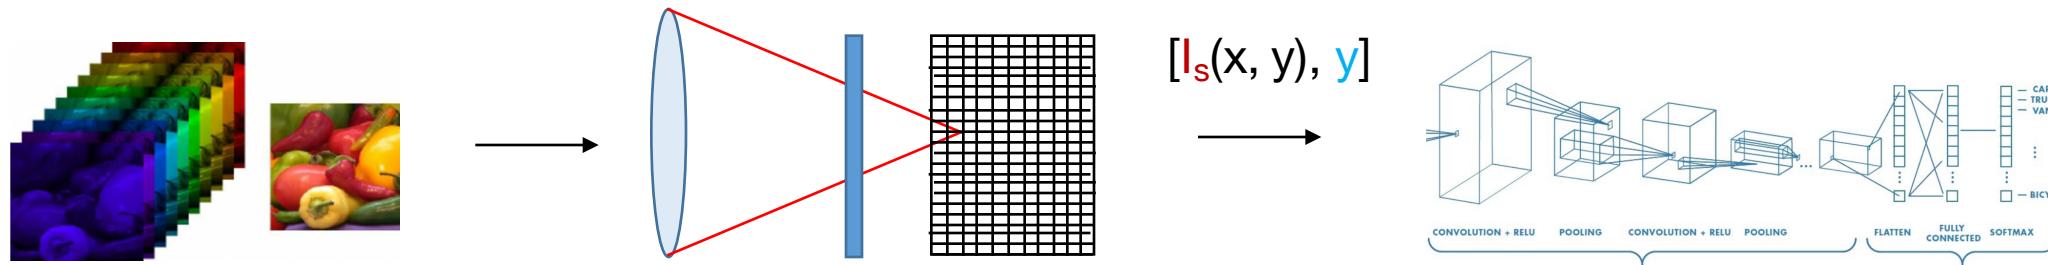
I_0 : 100 x 100 pix. x 30

Label y : 1x3 - pepper, broccoli, green beans

$$I_s(x, y) = \sum_{\lambda} W_0(\lambda) I_0(x, y, \lambda)$$

Physical Layer

Example #1: Optimized color filter for a grayscale camera



Training data:

$$[I_0(x, y, \lambda), y]$$

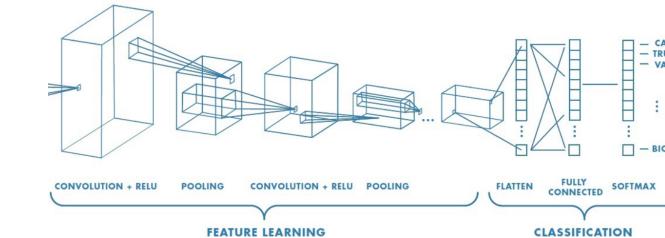
I_0 : $100 \times 100 \times 30$

Label y : 1×3

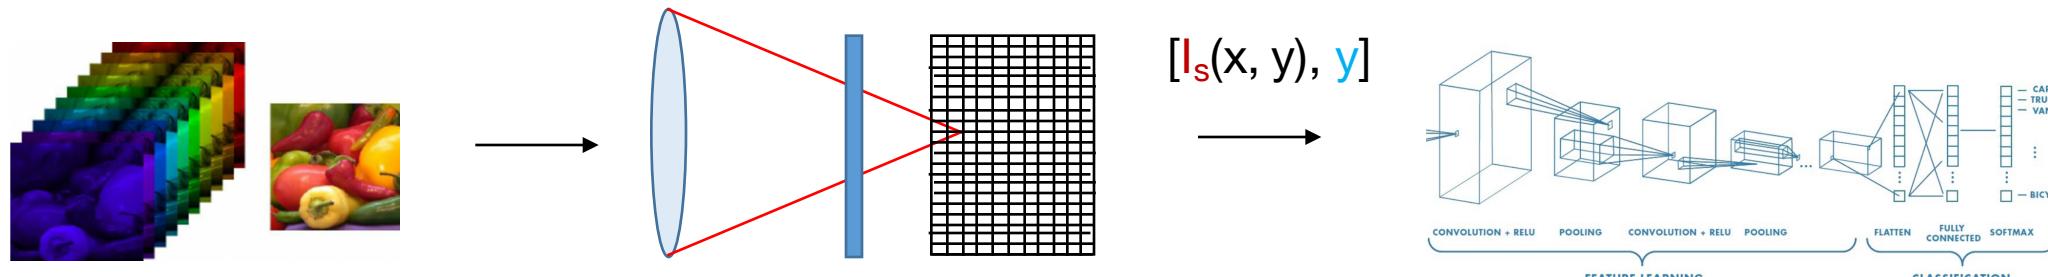
Physical Layer

$$I_s(x, y) = \sum_{\lambda} W_0(\lambda) I_0(x, y, \lambda)$$

$$\text{Task} = W_n \dots \text{ReLU}[W_1 \text{ReLU}[W_0 I_s] \dots]$$



Example #1: Optimized color filter for a grayscale camera



Training data:

$$[I_0(x, y, \lambda), y]$$

$I_0: 100 \times 100 \times 30$

Label $y: 1 \times 3$

$$I_s(x, y) = \sum_{\lambda} W_0(\lambda) I_0(x, y, \lambda)$$

Physical Layer

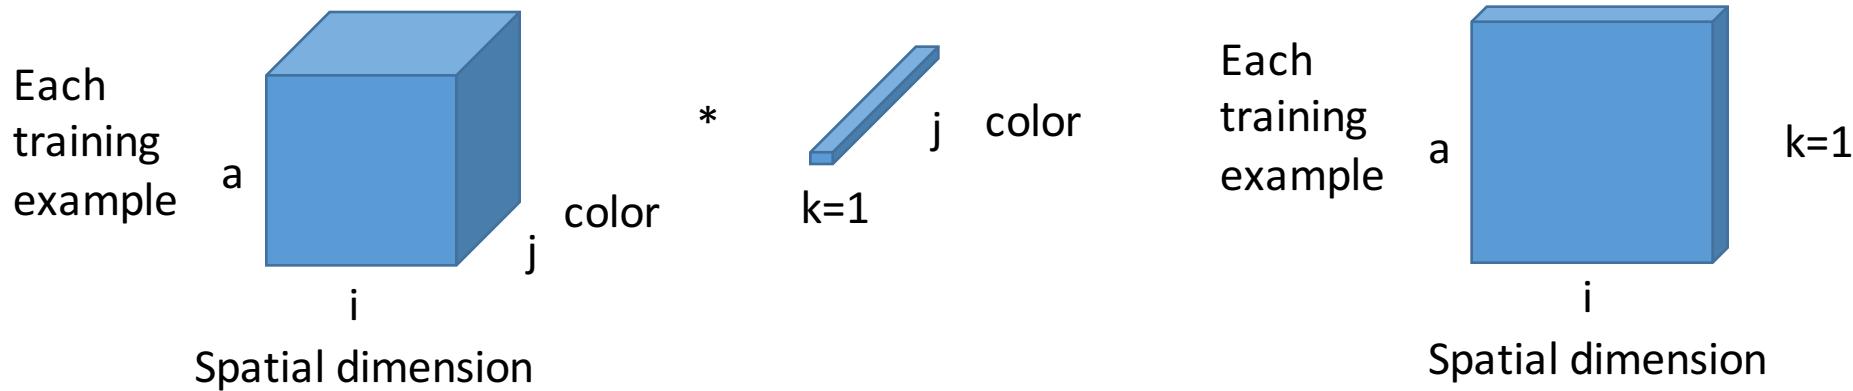
Task = $W_n \dots \text{ReLU}[W_1 \text{ReLU}[W_0 I_s] \dots]$

```
multispectral_data = tf.placeholder(tf.float32, [None, num_colors, image_size])
veg_labels = tf.placeholder(tf.float32, [None, 3])
filter_weights = tf.truncated_normal([num_colors, 1], stddev = 0.1)
filtered_images = tf.einsum('aij,jk->aij', multispectral_data, filter_weights)

#Now, train CNN of your choice using [filtered_images, veg_labels]
#Output of training/testing will be CNN weights, classification performance AND filter_weights!
```

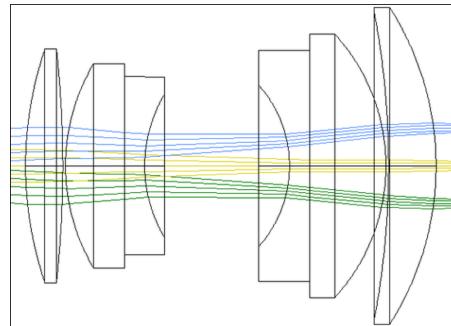
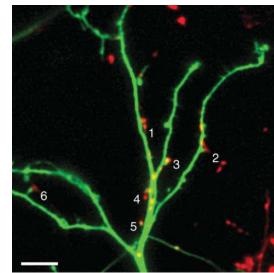
Tensorflow einsum – pretty useful for quick linear combo's

```
filtered_images = tf.einsum('aij,jk->aik', multispectral_data, filter_weights)
```



First model of light we'll consider – Incoherent light

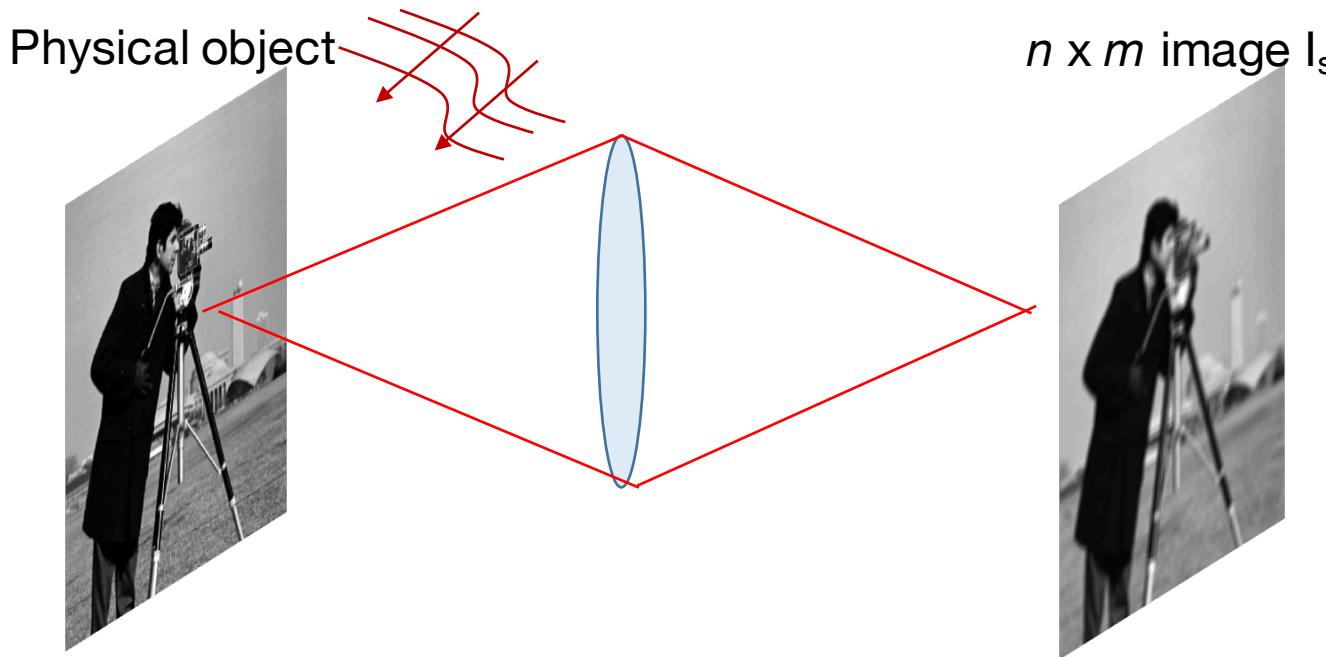
- Interpretation #1: Radiation (*Incoherent*)
- Model: Rays



- Real, non-negative
- Models absorption and brightness

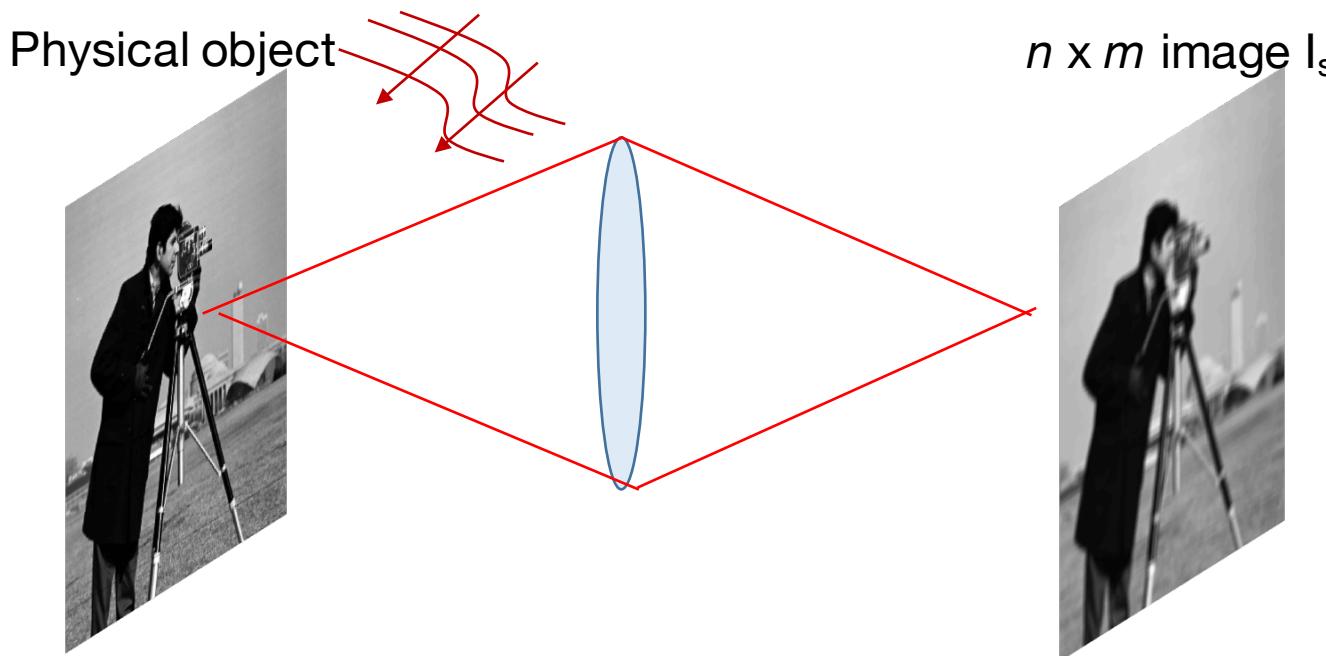
$$I_{\text{tot}} = I_1 + I_2$$

Example #2: Optimized illumination pattern (one color)



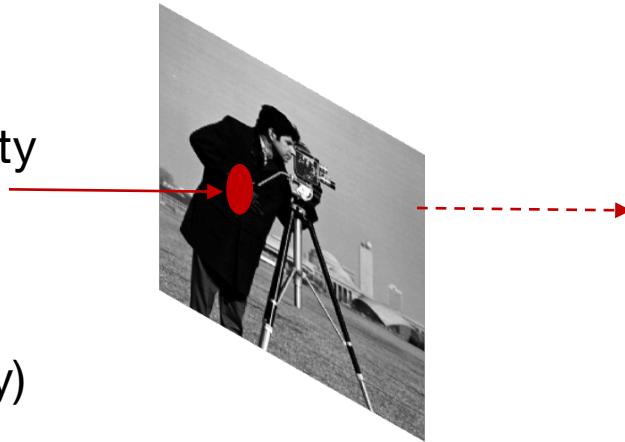
- Assume incoherent illumination
- Assume thin 2D object
- Object is real, non-negative map of absorption/reflectivity

Example #2: Optimized illumination pattern (one color)

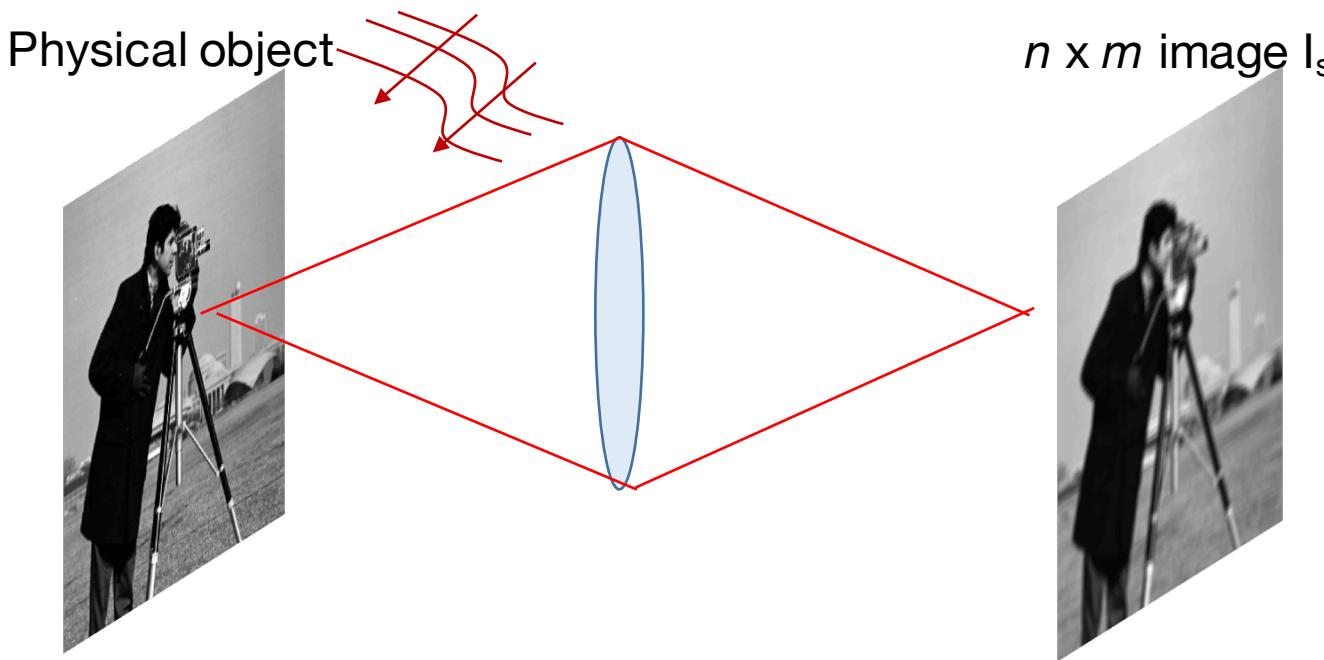


- Assume incoherent illumination
- Assume thin 2D object
- Object is real, non-negative map of absorption/reflectivity

Object absorption: $I_0(x,y)$
Illumination pattern: $s(x,y)$
Light exiting object surface: $I_e(x,y) = I_0(x,y) \circ s(x,y)$



Example #2: Optimized illumination pattern (one color)



- Assume incoherent illumination
- Assume thin 2D object
- Object is real, non-negative map of absorption/reflectivity

Modeling
incoherent
illumination

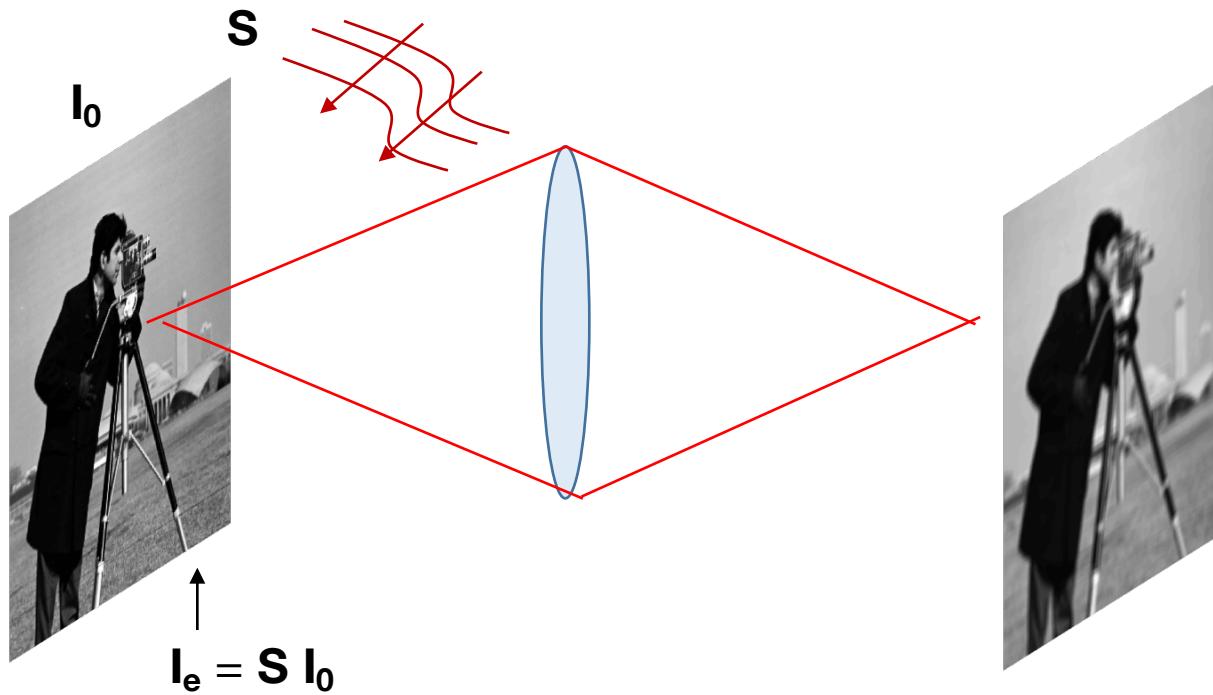
$$\mathbf{I}_e(x,y) = \mathbf{I}_0(x,y) \circ \mathbf{s}(x,y)$$

$$\mathbf{I}_e = \mathbf{S} \mathbf{I}_0$$

$$\text{diag}(\mathbf{S}) = \mathbf{s}$$

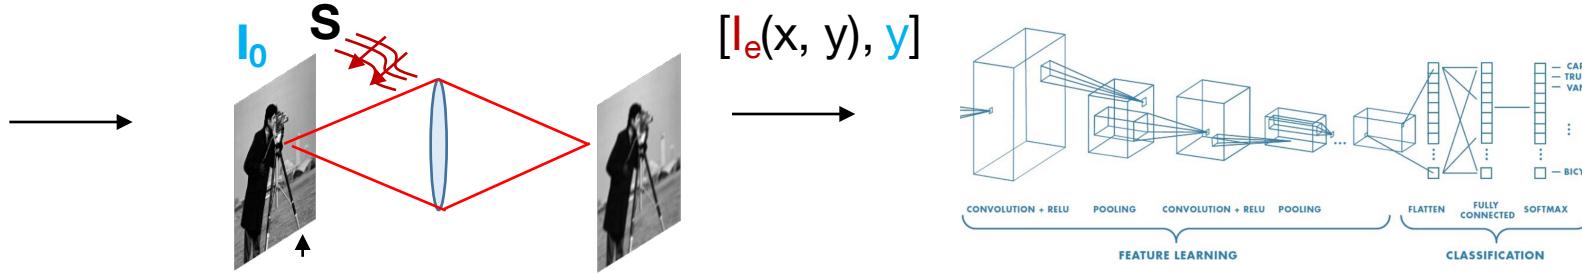
A square matrix \mathbf{S} is shown with a diagonal line running from the top-left corner to the bottom-right corner. The letter s is placed near the diagonal line.

Example #2: Optimized illumination pattern (one color)



First, assume perfect camera:
intensity at image plane $I_p = I_e = S I_0$

Example #2: Optimized illumination pattern (one color)

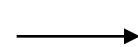


Training data:

$$[I_0(x, y), y]$$

I_0 : 100 x 100

Label y : 1x2

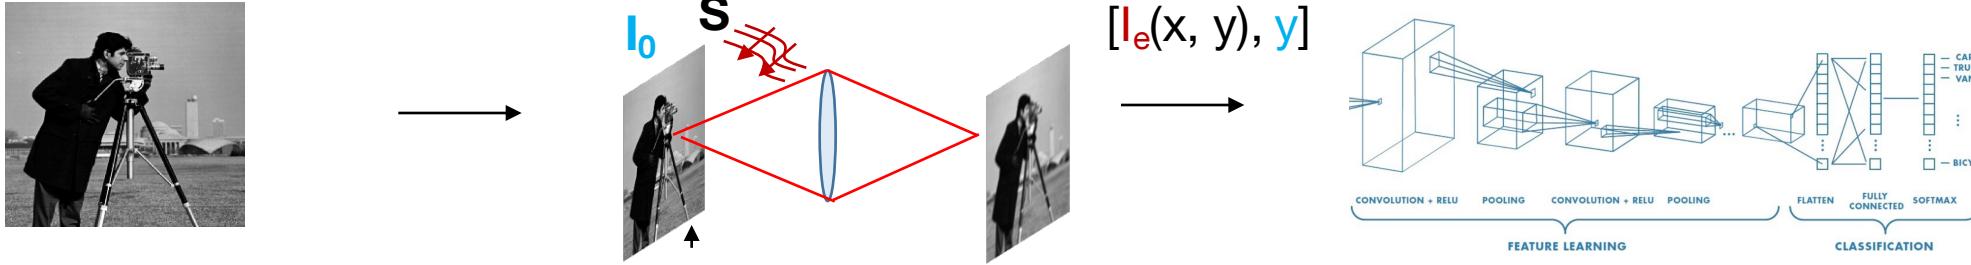


$$I_e(x, y) = \mathbf{S} I_0(x, y)$$

Physical Layer

$$\text{Task} = \mathbf{W}_n \dots \text{ReLU}[\mathbf{W}_1 \text{ReLU}[\mathbf{W}_0 I_e] \dots]$$

Example #2: Optimized illumination pattern (one color)



Training data:

$$[I_0(x, y), y]$$

$I_0: 100 \times 100$

Label $y: 1 \times 2$



$$I_e(x, y) = S I_0(x, y)$$

Physical Layer



$$\text{Task} = W_n \dots \text{ReLU}[W_1 \text{ReLU}[W_0 I_e] \dots]$$

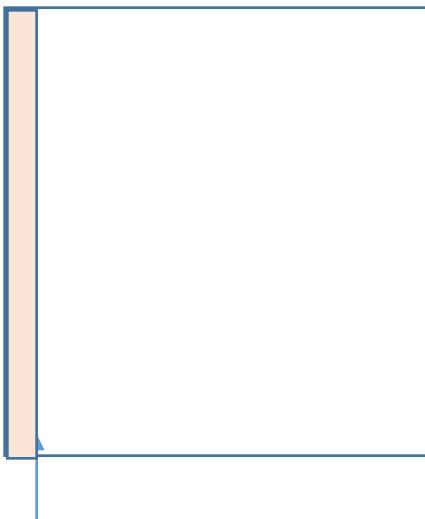
```
training_images = tf.placeholder(tf.float32, [image_size, None])
training_labels = tf.placeholder(tf.float32, [None, 3])
illumination_pattern = tf.truncated_normal([image_size, 1], stddev = 0.1)
illumination_matrix = tf.linalg.diag(illumination_pattern)
illumianted_images = tf.matmul(illumination_matrix, training_images)

#Now, train CNN of your choice using [illumianted_images, training_labels]
#Output of train/test will be CNN weights, classification performance AND illumination_pattern!
```

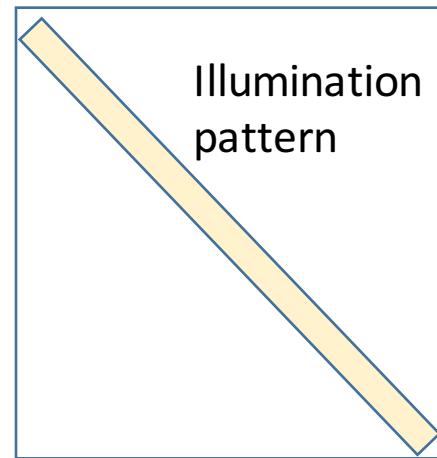
```
training_images = tf.placeholder(tf.float32, [image_size, None])
training_labels = tf.placeholder(tf.float32, [None, 3])
illumination_pattern = tf.truncated_normal([image_size, 1], stddev = 0.1)
illumination_matrix = tf.linalg.diag(illumination_pattern)
illumianted_images = tf.matmul(illumination_matrix, training_images)

#Now, train CNN of your choice using [illumianted_images, training_labels]
#Output of train/test will be CNN weights, classification performance AND illumination_pattern!
```

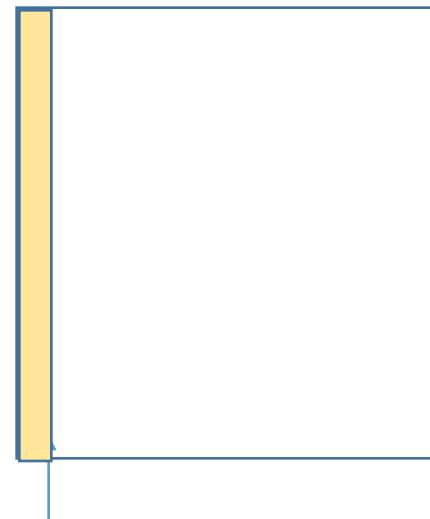
Illuminated images



Illumination Matrix



Training images



Illuminated Image 1

Image 1

We can also add in some lens blur

Lenses blur and rescale images:
(We'll learn how exactly next few weeks)

Input
intensity



$$\text{Input intensity} \quad * \quad \text{Convolution filter } h = \text{Output intensity}$$

The diagram illustrates the convolution process. On the left, labeled "Input intensity", is a sharp black and white photo of a man with a camera. In the center, labeled "Convolution filter h ", is a small, square, blurred kernel representing a lens blur effect. To the right, labeled "Output intensity", is a blurry version of the same photo, showing the effect of applying the filter to the input image.



We can also add in some lens blur

Lenses blur and rescale images:
(We'll learn how exactly next few weeks)

Input
intensity



$$\text{Convolution filter } h \\ * \quad \begin{matrix} & & \\ & & \\ & & \end{matrix}$$

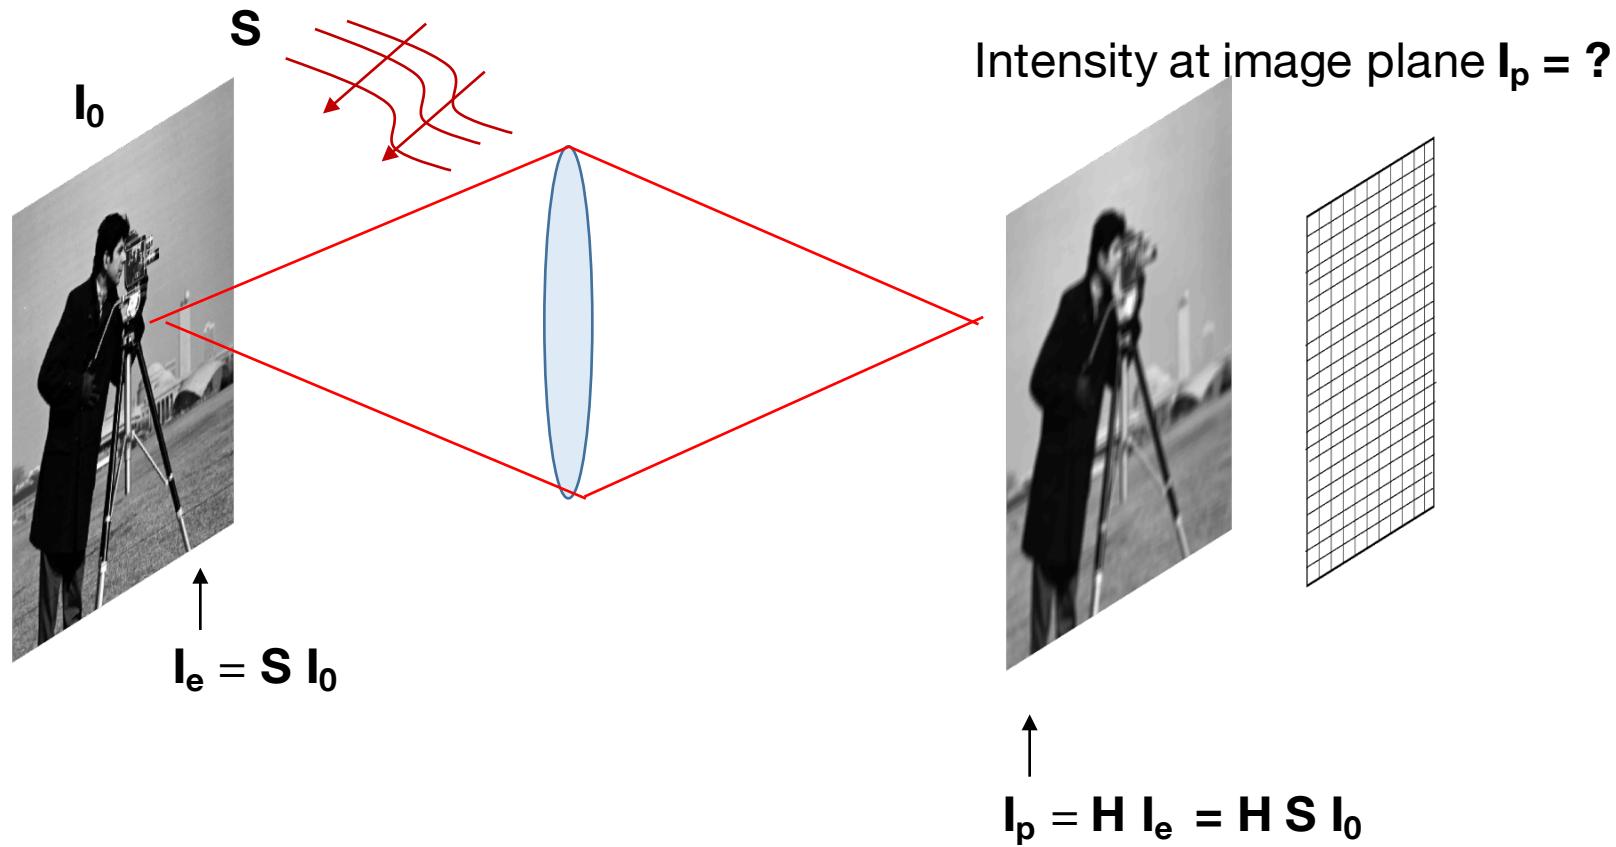


Output intensity

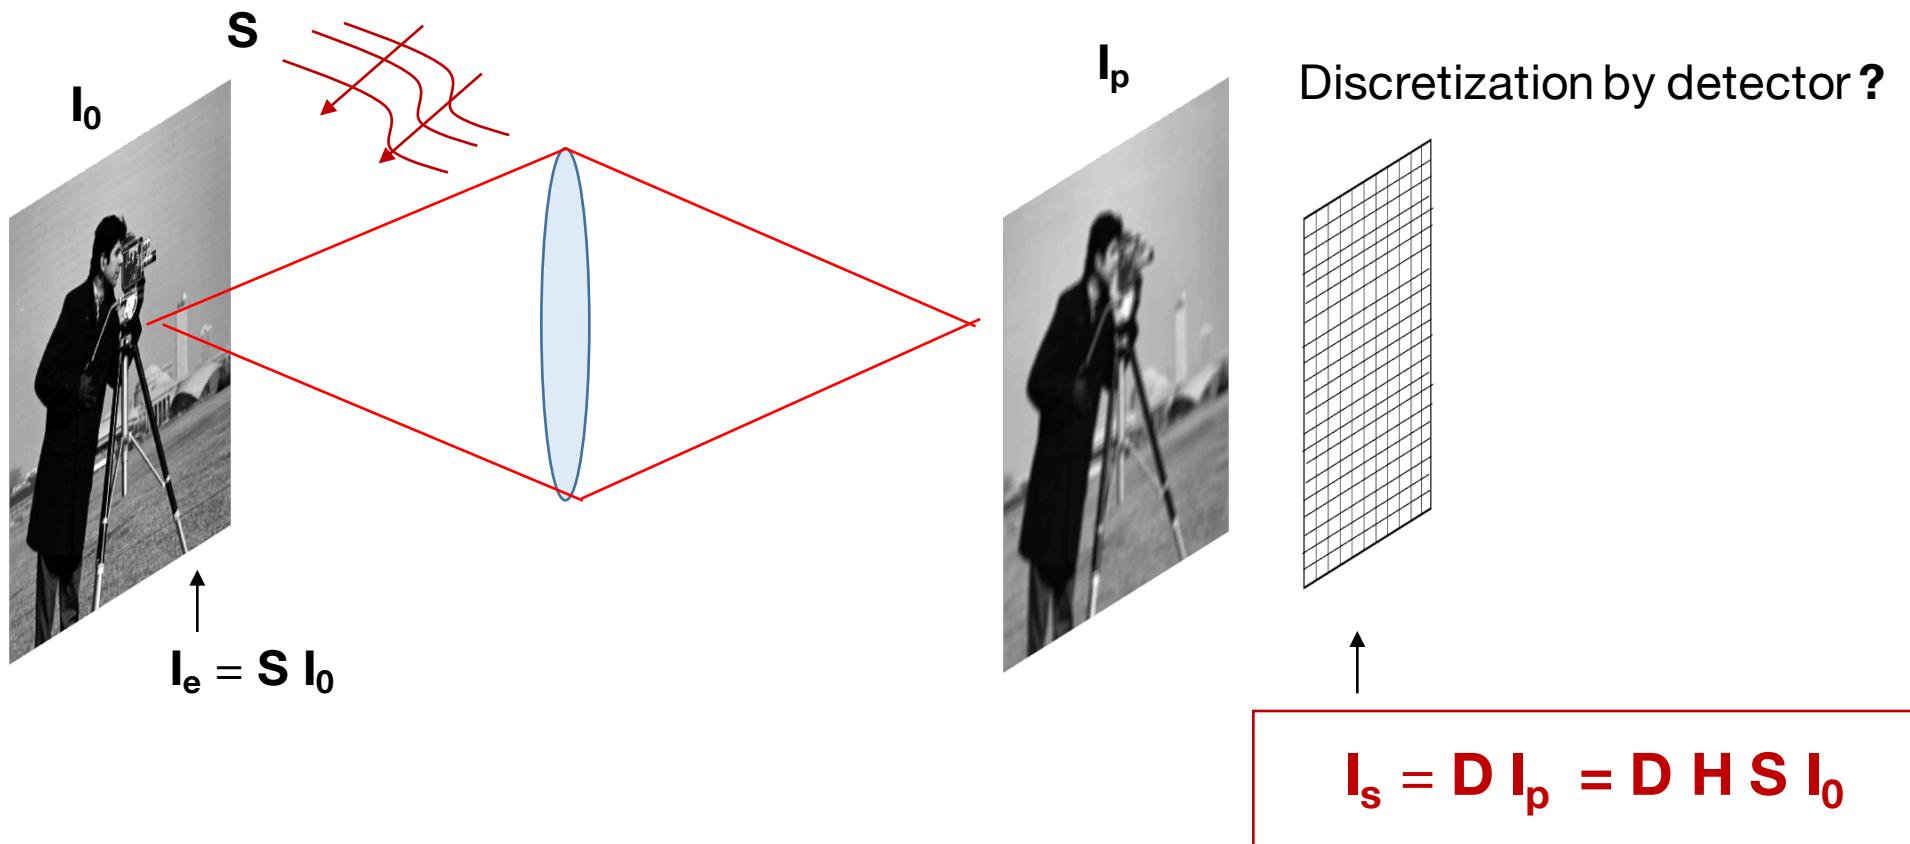
$$I_e(x/M, y/M) * h(x, y) = I_p(x, y)$$

Assuming we've resized by M, $I_p = I_e * h = H I_e$

Simple mathematical model of image formation



Simple mathematical model of image formation



Use downsampling matrix
(sum-pooling)

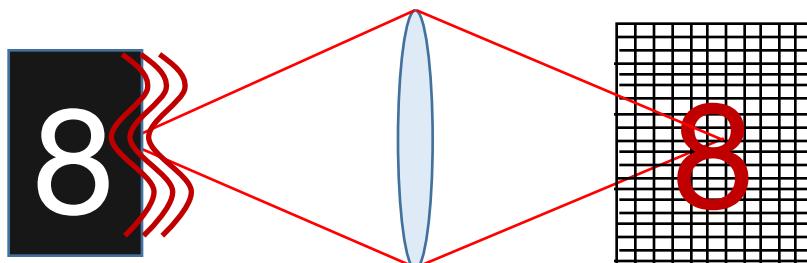
$$D =$$

0.5	0.5	0	0	0
0.5	0.5	0	0	0
0.5	0.5	0	0	0
⋮					⋮

Physical Layers

Physical world

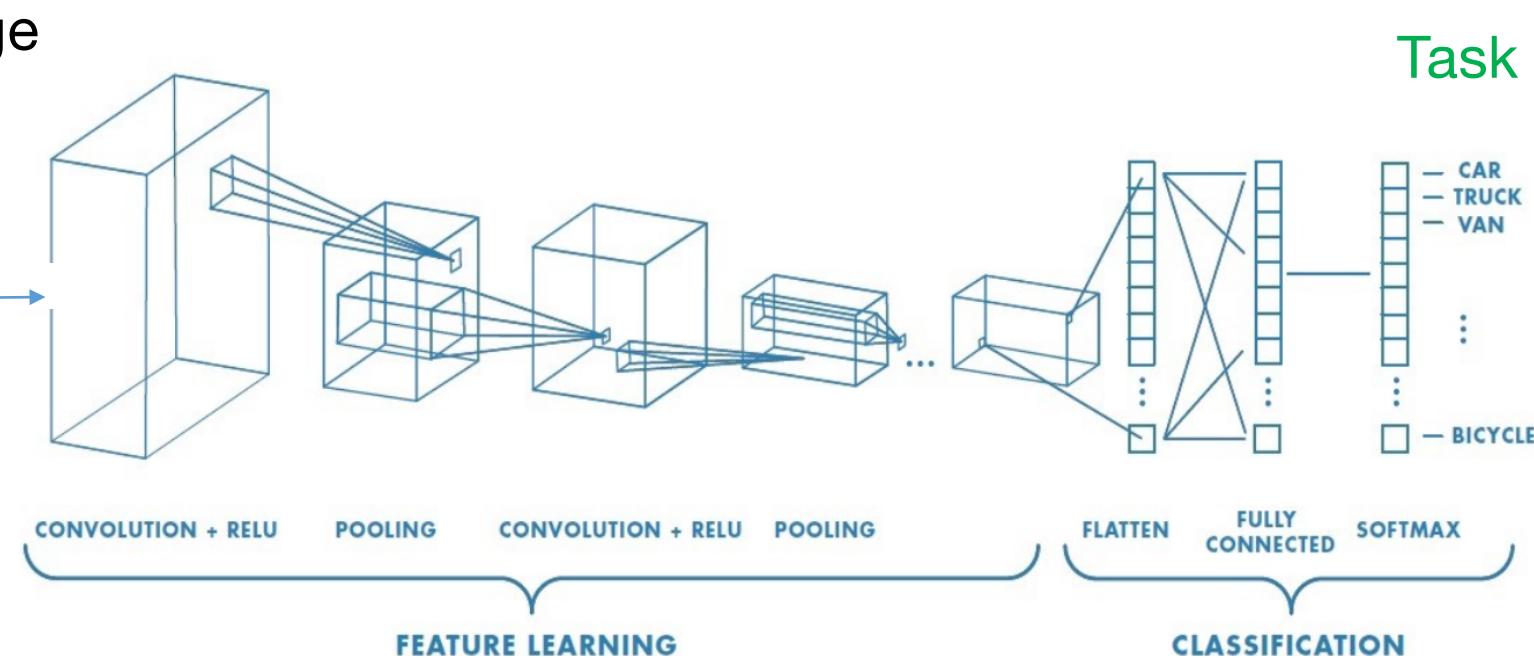
$$I_0(x_0, y_0)$$



Digital Image

$$I_s(x, y)$$

Digital Layers



Digital layers

$$\text{Task} = W_n \dots \text{ReLU}[W_1 \text{ReLU}[W_0 f[I_0]] \dots]$$

Physical layers

$$\text{Task} = W_n \dots \text{ReLU}[W_1 \text{ReLU}[W_0 \mathbf{D} \mathbf{H} \mathbf{S} I_0] \dots]$$

Summary: simple physical layers for incoherent imaging

- Deal with sample/image intensities I , real and non-negative
- Different colors add linearly

$$I_s(x, y) = \sum_{\lambda} I_0(x, y, \lambda)$$

- Effect of illumination is element-wise multiplication
- Imaging systems blur the object via point-spread function matrix H
- Discrete pixels down-sample the object via

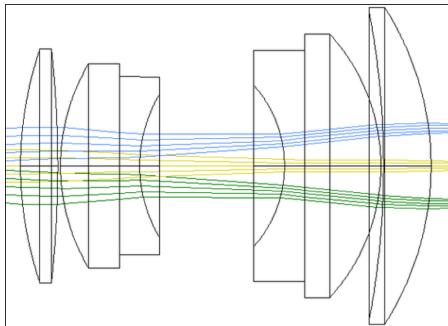
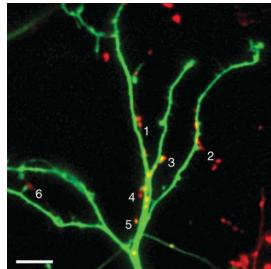
$$I_e(x, y) = S I_0(x, y)$$

$$I_b(x, y) = H I_0(x, y)$$

$$I_d(x, y) = D I_0(x, y)$$

First - what is light and how can we model it?

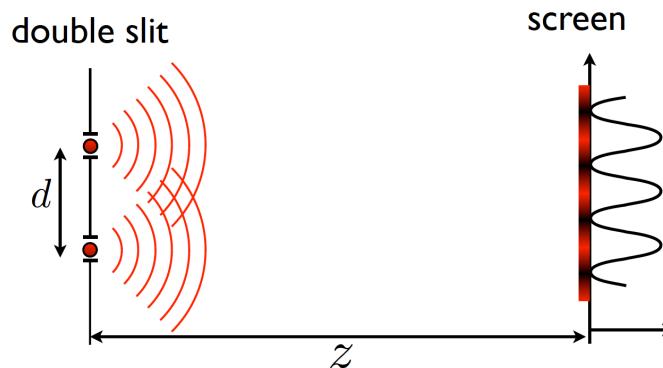
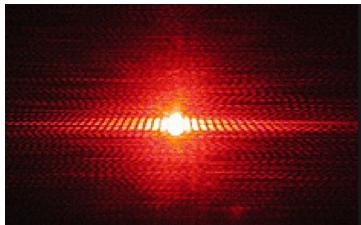
- Interpretation #1: Radiation (*Incoherent*)
- Model: Rays



- Real, non-negative
- Models absorption and brightness

$$I_{\text{tot}} = I_1 + I_2$$

- Interpretation #2: Electromagnetic wave (*Coherent*)
- Model: Waves



- Complex field
- Models Interference

$$E_{\text{tot}} = E_1 + E_2$$

Let's take a step back: how does light propagate?

Maxwell's equations
without any charge

$$\nabla \times \vec{\mathcal{E}} = -\mu \frac{\partial \vec{\mathcal{H}}}{\partial t}$$

$$\nabla \times \vec{\mathcal{H}} = \epsilon \frac{\partial \vec{\mathcal{E}}}{\partial t}$$

$$\nabla \cdot \epsilon \vec{\mathcal{E}} = 0$$

$$\nabla \cdot \mu \vec{\mathcal{H}} = 0.$$

Let's take a step back: how does light propagate?

Maxwell's equations
without any charge

$$\nabla \times \vec{\mathcal{E}} = -\mu \frac{\partial \vec{\mathcal{H}}}{\partial t}$$

$$\nabla \times \vec{\mathcal{H}} = \epsilon \frac{\partial \vec{\mathcal{E}}}{\partial t}$$

$$\nabla \cdot \epsilon \vec{\mathcal{E}} = 0$$

$$\nabla \cdot \mu \vec{\mathcal{H}} = 0.$$

1. Take the curl of both sides of first equation
2. Substitute 2nd and 3rd equation
3. Arrive at the wave equation:

$$\nabla^2 \vec{\mathcal{E}} - \frac{n^2}{c^2} \frac{\partial^2 \vec{\mathcal{E}}}{\partial t^2} = 0 \quad n = \left(\frac{\epsilon}{\epsilon_0} \right)^{1/2} \quad c = \frac{1}{\sqrt{\mu_0 \epsilon_0}}.$$

Let's take a step back: how does light propagate?

Considering light that isn't pulsed over time, we can use the following solution:

$$u(P, t) = A(P) \cos[2\pi\nu t + \phi(P)]$$

$$u(P, t) = \text{Re}\{U(P) \exp(-j2\pi\nu t)\},$$

Let's take a step back: how does light propagate?

Considering light that isn't pulsed over time, we can use the following solution:

$$u(P, t) = A(P) \cos[2\pi\nu t + \phi(P)]$$

$$u(P, t) = \text{Re}\{U(P) \exp(-j2\pi\nu t)\},$$

With this particular solution, we get the following important time-independent equation:

Helmholtz
Equation

$$(\nabla^2 + k^2)U = 0.$$

$$k = 2\pi n \frac{\nu}{c} = \frac{2\pi}{\lambda},$$

Let's take a step back: how does light propagate?

Considering light that isn't pulsed over time, we can use the following solution:

$$u(P, t) = A(P) \cos[2\pi\nu t + \phi(P)]$$

$$u(P, t) = \text{Re}\{U(P) \exp(-j2\pi\nu t)\},$$

With this particular solution, we get the following important time-independent equation:

Helmholtz
Equation

$$(\nabla^2 + k^2)U = 0.$$

$$k = 2\pi n \frac{\nu}{c} = \frac{2\pi}{\lambda},$$

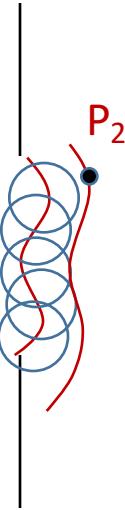
This is an important equation in physics. We won't go into the details, but it leads to the Huygen-Fresnel principle:

$$U(P_2) = \frac{1}{j\lambda} \iint_{\Sigma} U(P_1) \frac{\exp(jkr_{21})}{r_{21}} \cos\theta ds$$

Plane-to-plane light propagation via the "paraxial approximation"

The Huygens-Fresnel Equation

$$U(P_2) = \frac{1}{j\lambda} \iint_{\Sigma} U(P_1) \frac{\exp(jkr_{21})}{r_{21}} \cos \theta ds$$

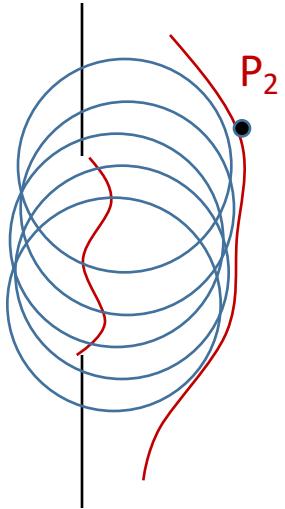


Aperture

Plane-to-plane light propagation via the "paraxial approximation"

The Huygens-Fresnel Equation

$$U(P_2) = \frac{1}{j\lambda} \iint_{\Sigma} U(P_1) \frac{\exp(jkr_{21})}{r_{21}} \cos \theta ds$$



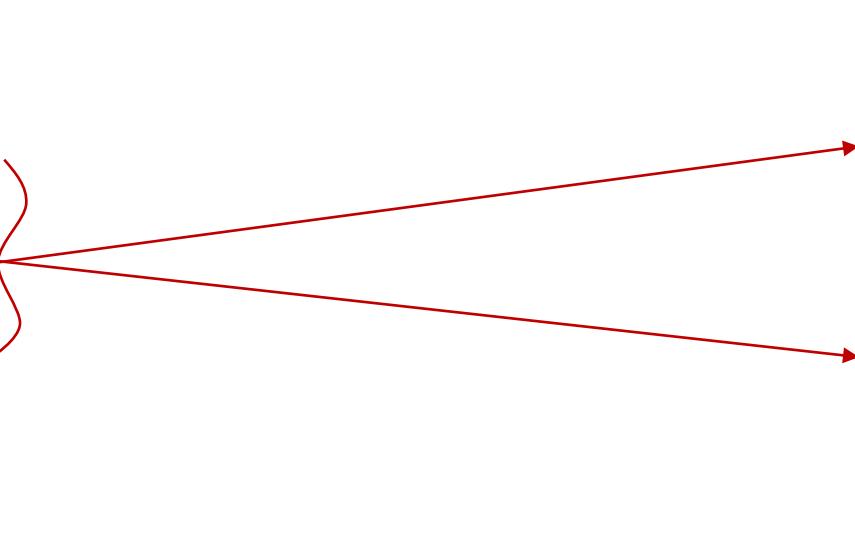
Generally connects two points in 3D:

$$U(P_1) = U(x_1, y_1, z_1)$$

$$U(P_2) = U(x_2, y_2, z_2)$$

Plane-to-plane light propagation via the "paraxial approximation"

We are usually concerned about propagation between two planes (almost always in an optical system):



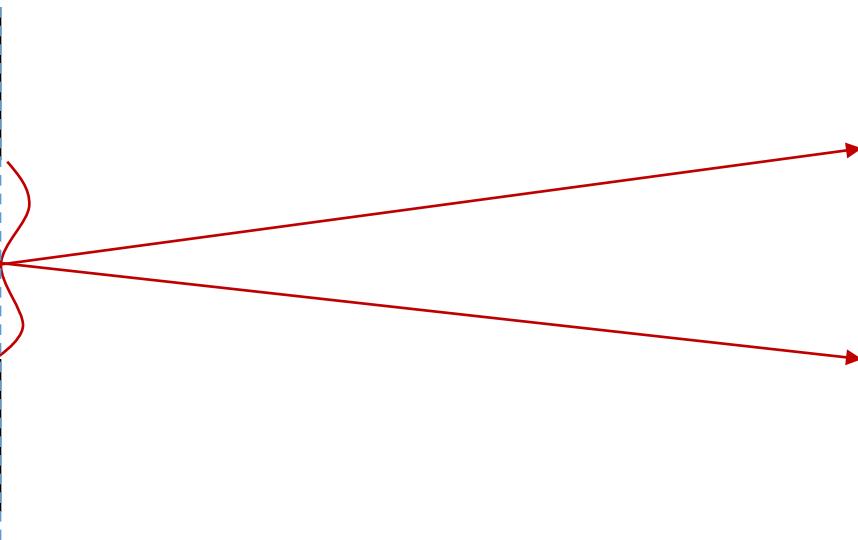
$$U(P_1) = U(x_1, y_1, z_1 = z_{p1})$$

$$U(P_2) = U(x_2, y_2, z_2 = z_{p2})$$

$$U(P) = E(x, y, z) e^{ikz}$$

Plane-to-plane light propagation via the "paraxial approximation"

We are usually concerned about propagation between two planes (almost always in an optical system):



Paraxial approximation:

$$\nabla_{\perp}^2 U + 2ik \frac{dU}{dz} = 0$$

$$\nabla_{\perp}^2 \stackrel{\text{def}}{=} \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$

$$U(P_1) = U(x_1, y_1, z_1 = z_{p1})$$

$$U(P_2) = U(x_2, y_2, z_2 = z_{p2})$$

$$U(P) = E(x, y, z) e^{ikz}$$

Plane-to-plane light propagation via the "paraxial approximation"

We are usually concerned about propagation between two planes (almost always in an optical system):

Paraxial approximation:

$$\nabla_{\perp}^2 U + 2ik \frac{dU}{dz} = 0 \quad \text{Substitute in } U(P) = E(x, y, z)e^{ikz} \text{ and crank the wheel,}$$

$$\nabla_{\perp}^2 E + 2ik \frac{dE}{dz} + 2k^2 E = 0 \quad \text{Paraxial Helmholtz Equation. This has an exact integral solution:}$$

Plane-to-plane light propagation via the "paraxial approximation"

We are usually concerned about propagation between two planes (almost always in an optical system):

Paraxial approximation:

$$\nabla_{\perp}^2 U + 2ik \frac{dU}{dz} = 0 \quad \text{Substitute in } U(P) = E(x, y, z)e^{ikz} \text{ and crank the wheel,}$$

$$\nabla_{\perp}^2 E + 2ik \frac{dE}{dz} + 2k^2 E = 0 \quad \text{Paraxial Helmholtz Equation. This has an exact integral solution:}$$

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint_{-\infty}^{+\infty} E(x', y', 0) e^{\frac{ik}{2z} [(x-x')^2 + (y-y')^2]} dx' dy'$$

Fresnel diffraction
integral

This is how light propagates from one plane to the next. It's a convolution!

Fresnel light propagation as a convolution

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint_{-\infty}^{+\infty} E(x', y', 0) e^{\frac{ik}{2z} [(x-x')^2 + (y-y')^2]} dx' dy'$$

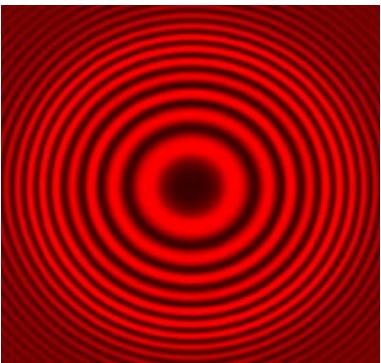
$$h(x, y, z) = \frac{e^{ikz}}{i\lambda z} e^{i\frac{k}{2z}(x^2 + y^2)}$$

$$E(x, y, z) = E(x, y, 0) * h(x, y, z)$$

Fresnel light propagation as a convolution

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint_{-\infty}^{+\infty} E(x', y', 0) e^{\frac{ik}{2z} [(x-x')^2 + (y-y')^2]} dx' dy'$$

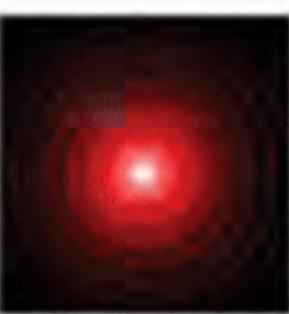
$$h(x, y, z) = \frac{e^{ikz}}{i\lambda z} e^{i\frac{k}{2z}(x^2 + y^2)}$$



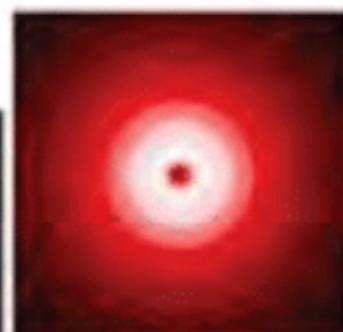
$$E(x, y, z) = E(x, y, 0) * h(x, y, z)$$

Paraxial
image
plane

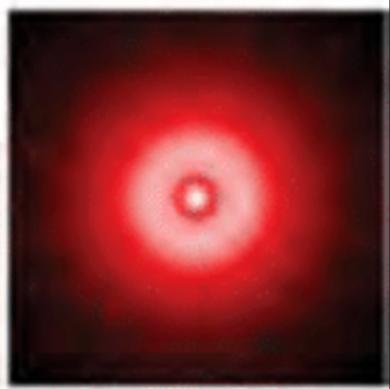
$W_{020} = 0$



$W_{020} = -\lambda/2$



$W_{020} = -\lambda$



$W_{020} = -3\lambda/2$

From the Fresnel approximation to the Fraunhofer approximation

Fresnel Approximation:

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint_{-\infty}^{+\infty} E(x', y', 0) e^{\frac{ik}{2z} [(x-x')^2 + (y-y')^2]} dx' dy'$$

Lets assume that the second plane is “pretty far away” from the first plane. Then,

$$z > \frac{2D^2}{\lambda}$$

From the Fresnel approximation to the Fraunhofer approximation

Fresnel Approximation:

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint_{-\infty}^{+\infty} E(x', y', 0) e^{\frac{ik}{2z} [(x-x')^2 + (y-y')^2]} dx' dy'$$

Lets assume that the second plane is “pretty far away” from the first plane. Then,

$$z > \frac{2D^2}{\lambda}$$

1. Expand the squaring

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint E(x', y', 0) e^{\frac{ik}{2z} (x^2 + y^2)} e^{\frac{ik}{2z} (x'^2 + y'^2)} e^{\frac{ik}{2z} (xx' + yy')} dx' dy'$$

From the Fresnel approximation to the Fraunhofer approximation

Fresnel Approximation:

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint_{-\infty}^{+\infty} E(x', y', 0) e^{\frac{ik}{2z} [(x-x')^2 + (y-y')^2]} dx' dy'$$

Lets assume that the second plane is “pretty far away” from the first plane. Then,

$$z > \frac{2D^2}{\lambda}$$

1. Expand the squaring

$$E(x, y, z) = \frac{e^{ikz}}{i\lambda z} \iint E(x', y', 0) e^{\frac{ik}{2z} (x^2 + y^2)} e^{\frac{ik}{2z} (x'^2 + y'^2)} e^{\frac{ik}{2z} (xx' + yy')} dx' dy'$$

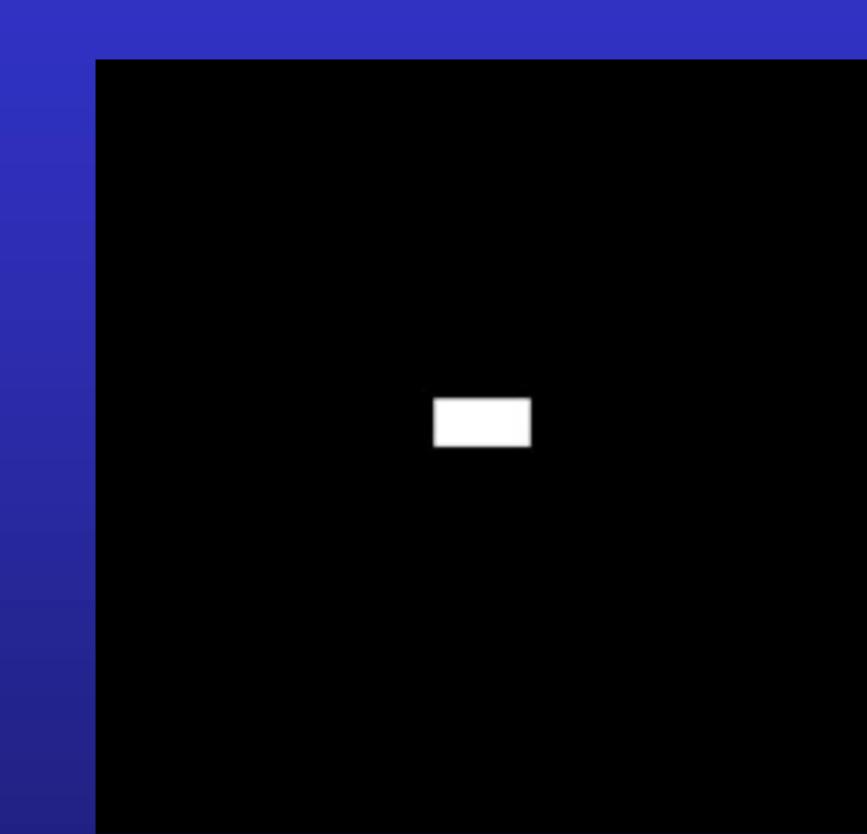
2. Front term comes out, assume second term goes away, then,

$$E(x, y, z) = C \iint E(x', y', 0) e^{\frac{ik}{2z} (xx' + yy')} dx' dy'$$

$$C = \frac{e^{ikz}}{i\lambda z} e^{\frac{ik}{2z} (x^2 + y^2)}$$

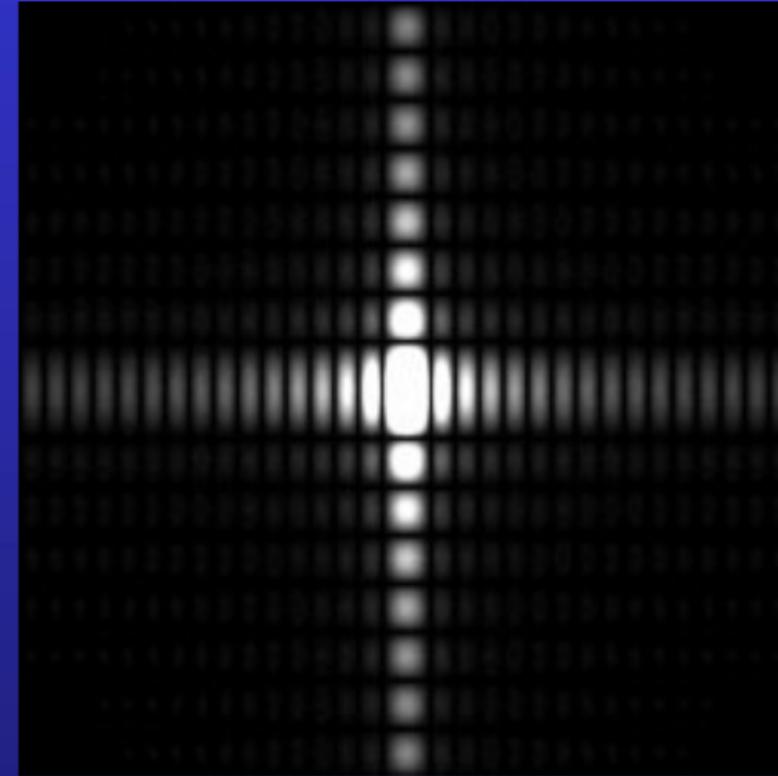
Fraunhofer diffraction is a Fourier transform!!!!!!

This is the aperture



Two-dimensional rectangle
function as an image

This is what you see far away



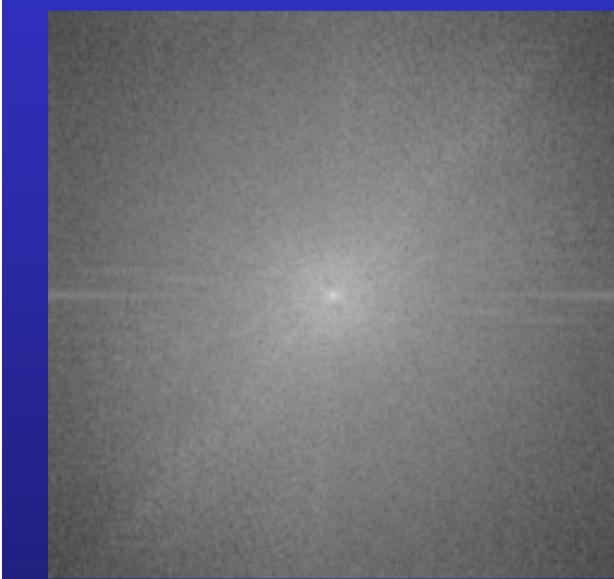
d) Magnitude of Fourier spectrum
of the 2-D rectangle



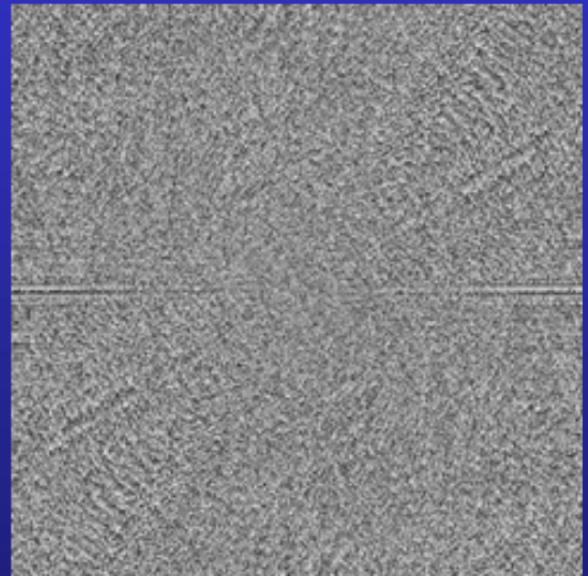
Cheetah



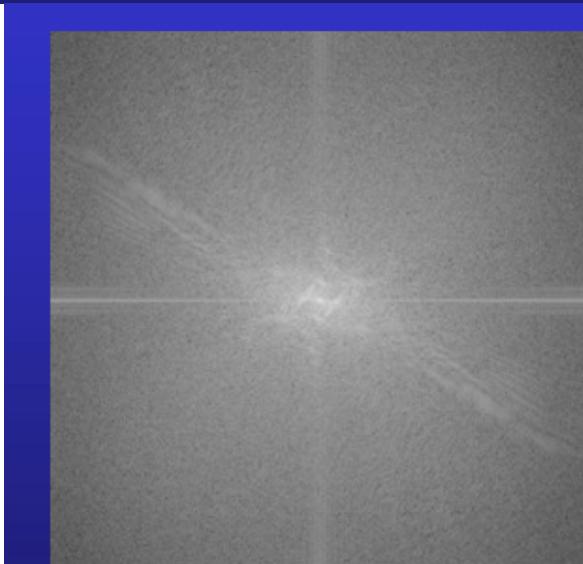
Zebra



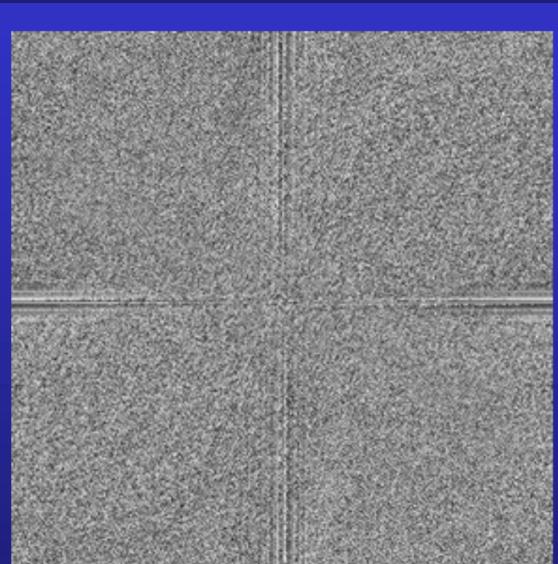
magnitude of cheetah



phase of cheetah



magnitude of zebra



phase of zebra